

Introduction to PCA, k-means clustering, and k nearest neighbors classification

Yulia Newton, Ph.D.



CS156, Introduction to Artificial Intelligence

San Jose State University

Spring 2021

Data dimensionality

- “Dimensionality” refers to the number of features each training/validation/testing observation has




Before you jump into building predictive models

- Not all of the features in your dataset might be informative or relevant to the prediction task you are solving
 - Irrelevant information adds unnecessary noise
 - Reducing dimensionality can help improve model performance
 - Noise can cause model overfitting
- The curse of dimensionality
 - Models built on high-dimensional data require a lot of observations in order to avoid overfitting
 - Reducing dimensionality can help avoid model overfitting
- Remember: “garbage in - garbage out”

Dimensionality reduction

- “Dimensionality reduction, or dimension reduction, is the transformation of data from a high-dimensional space into a low-dimensional space so that the low-dimensional representation retains some meaningful properties of the original data, ideally close to its intrinsic dimension.” - Wikipedia



Feature space reduction


- Dimensionality reduction is also known as a feature space reduction
- Can be achieved through
 - Reducing the number of input/independent variables
 - Transformation/projection/rotation of the data from high-dimensional feature space into low-dimensional feature space

Dimensionality reduction during data exploration stage

- Projecting data into 2-D or 3-D space prior to building a predictive model can show patterns intrinsic to the data
- Clustering the data prior to building a predictive model can tell us if the chosen feature space is predictive of the output variable

Principal component analysis


- Principal component analysis (PCA) is one of the most commonly used methods to perform data projection and dimensionality reduction
- Core idea - rotate axes in the direction of the biggest variance



<https://amva4newphysics.wordpress.com>

PCA steps (conceptual understanding)

- Find the direction in the data with the highest variance in the hyperspace (feature space) - that is the first principal component
- Find the direction in the data with the next highest variance, which is orthogonal to previously identified principal components
- Continue until the number of desired principal components is reached



<https://www.analyticsvidhya.com/blog/2016/03/pca-practical-guide-principal-component-analysis-python>


What are principal components?

- Principal components are eigenvectors of the covariance matrix
 - From linear algebra, an eigenvector is a vector that changes only by a scalar under a linear transformation
 - Does not change direction in a transformation

$$A\mathbf{v} = \lambda\mathbf{v}$$

Diagram illustrating the equation $A\mathbf{v} = \lambda\mathbf{v}$. Arrows point from the labels below to the corresponding terms in the equation: "Matrix" points to A , "Eigenvector" points to \mathbf{v} on the left, "Eigenvalue" points to λ , and "Eigenvector" points to \mathbf{v} on the right.

www.mathsisfun.com



Mathematics behind PCA

- Compute covariance matrix of the input data
 - X_1, X_2, \dots, X_n are random variables
- Covariance matrix is a matrix of covariances between random variables i and j in each (i, j) entry of the matrix
 - On the diagonal, it's just variance of the random variable
 - According to the law of large numbers, the average value of a random variable converges to expected value

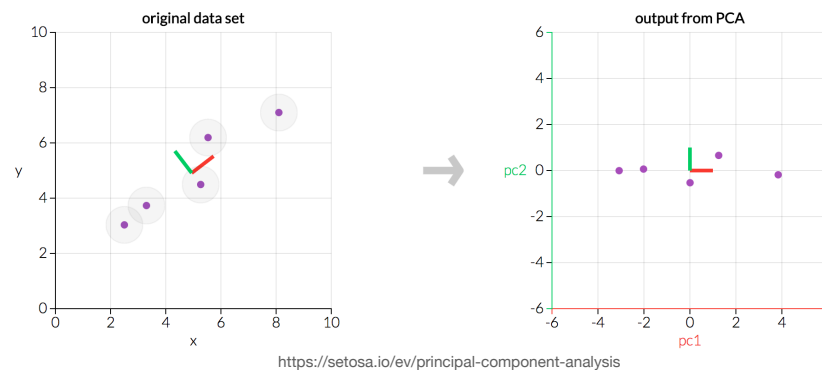
$$\mathbf{X} = (X_1, X_2, \dots, X_n)^T$$

$$K_{X_i X_j} = \text{cov}[X_i, X_j] = E[(X_i - E[X_i])(X_j - E[X_j])]$$

$$K_{\mathbf{X}\mathbf{X}} = \begin{bmatrix} E[(X_1 - E[X_1])(X_1 - E[X_1])] & E[(X_1 - E[X_1])(X_2 - E[X_2])] & \cdots & E[(X_1 - E[X_1])(X_n - E[X_n])] \\ E[(X_2 - E[X_2])(X_1 - E[X_1])] & E[(X_2 - E[X_2])(X_2 - E[X_2])] & \cdots & E[(X_2 - E[X_2])(X_n - E[X_n])] \\ \vdots & \vdots & \ddots & \vdots \\ E[(X_n - E[X_n])(X_1 - E[X_1])] & E[(X_n - E[X_n])(X_2 - E[X_2])] & \cdots & E[(X_n - E[X_n])(X_n - E[X_n])] \end{bmatrix}$$

Principal components are not the same as the original features

- Once the data has been rotated into the principal component space, original features no longer can be assigned to the PC axes
 - Principal components are linear combinations of original features
 - Different original features might contribute to principal components with different influence
 - Principal component space is a new transformed space




Let's play with some online examples

- <https://setosa.io/ev/principal-component-analysis/>


Clustering

- “Cluster analysis or clustering is the task of grouping a set of objects in such a way that objects in the same group are more similar to each other than to those in other groups.” - Wikipedia
- Clustering assigns labels based on the input variable values



k-mean clustering

- This method aims to partition data into k parts/clusters
- The idea is to represent clusters by centroids and assign a data point to a cluster based on which centroid it is the closest to/most similar to




<https://towardsdatascience.com>

k-mean clustering (cont'd)

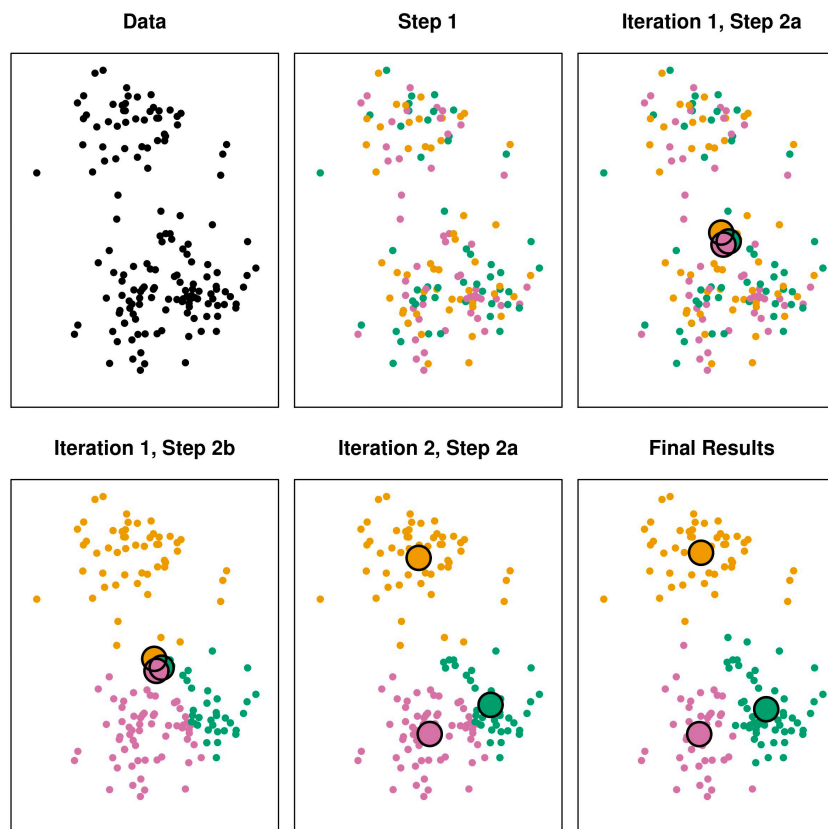
- Steps:
 - Select k parameter
 - Select k centroids by selecting random data points from the dataset without replacement
 - Other ways to generate initial centroids can be utilized
 - Iterate until convergence (no significant change in centroids and/or assignment of data points to the clusters does not change or a predetermined number of iterations):
 - For each data point compute distance or similarity between the data point and each centroid
 - Assign each data point to 1 out of k clusters based on the smallest distance/largest similarity
 - Recompute the centroid for each cluster 1...k based on the data points in that cluster

Distance/similarity measure

- Depending on the data/task different measures of distance/similarity might be more appropriate
 - Distance similarity can be measured in low- or high-dimensional space
- Most often used measures of distance:
 - Euclidean distance
 - Hamming distance
 - Manhattan distance
 - Minkowski distance
- Most often used measures of similarity:
 - Correlation coefficient
 - Cosine similarity
 - Jaccard similarity
- Useful explanation of some of these measures:
 - <https://medium.com/@gshriya195/top-5-distance-similarity-measures-implementation-in-machine-learning-1f68b9ecb0a3>

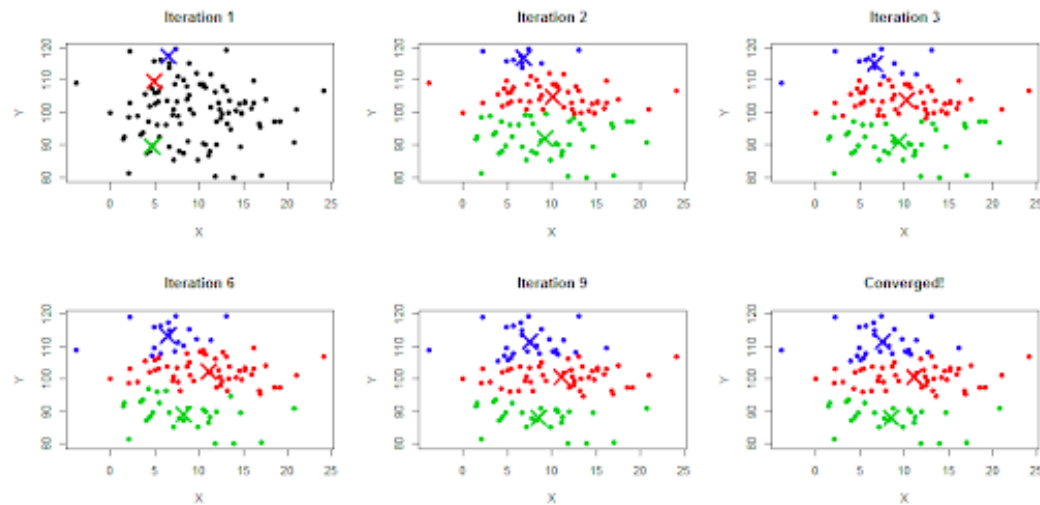


Demonstration of k-means clustering steps (toy example 1)



<https://stackoverflow.com/questions/51263331/kmeans-save-each-iteration-step>

Demonstration of k-means clustering steps (toy example 2)



<https://stackoverflow.com/questions/60312401/when-using-the-k-means-clustering-algorithm-is-it-possible-to-have-a-set-of-dat>

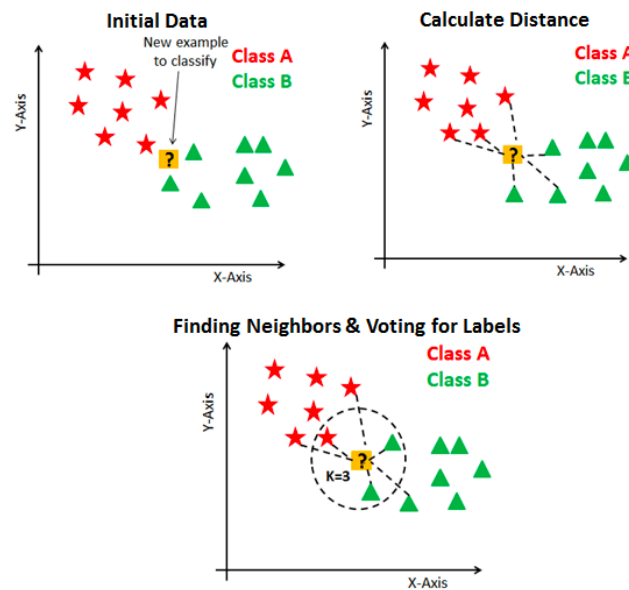
k-nearest neighbors classification

- k-nearest neighbors algorithm is a ML classification method for classifying observations into classes/labels
 - The output is a class membership
- No generalized model
 - Based on training data and labels
 - No explicit model training step
- An instance-based learning type of algorithm
 - Also called lazy learning
- Related to the k-means clustering algorithm

k-nearest neighbors method


- Steps for classifying a new data point:
 - Select k parameter
 - Compute the distance/similarity between every data point in the training data and the new data point
 - Identify k closest (most similar) data points
 - Assign the label of the new data point based on majority vote
 - Different ways of assigning the class can be utilized
 - E.g. weighted voting, etc.

k-nearest neighbors toy example



<https://www.datacamp.com/community/tutorials/k-nearest-neighbor-classification-scikit-learn>

The choice of k hyper-parameter might dramatically affect the results



Let's look at some code examples

- General dimensionality reduction examples on Iris dataset
 - scikitlearn libraries
 - PCA
 - LDA
 - t-SNE
 - *Dimensionality_reduction.ipynb*
- k-means clustering of synthetic 2-D data
 - *kmeans.synthetic_data.ipynb*
- k-nearest neighbors classification
 - *knn.synthetic_data.ipynb*