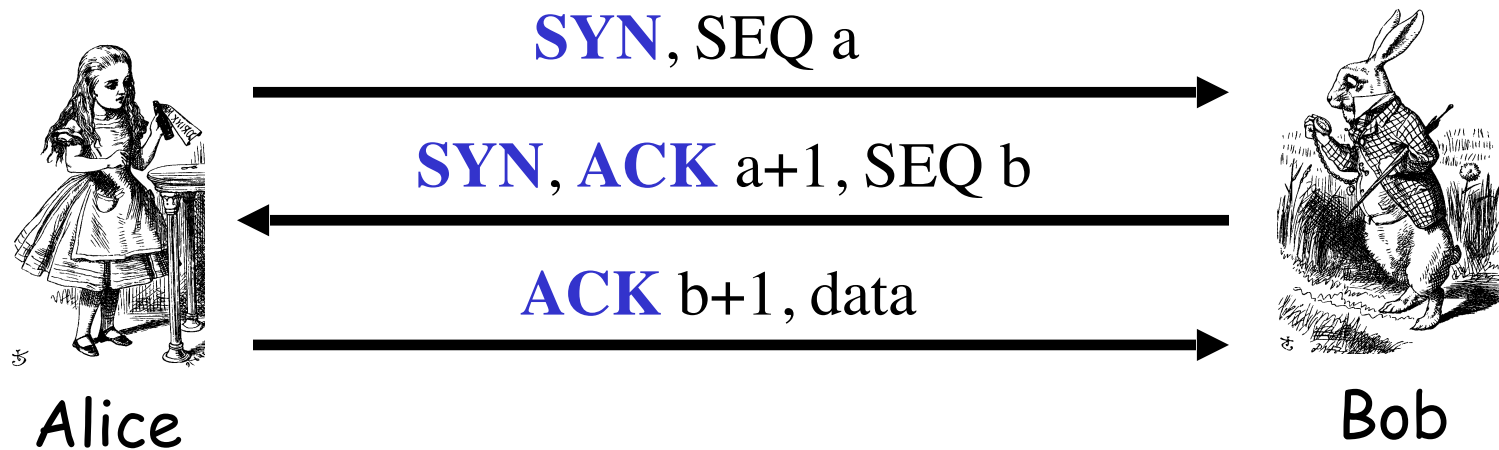


Authentication and TCP

TCP-based Authentication


- ❑ TCP not intended for use as an authentication protocol
- ❑ But IP address in TCP connection may be (mis)used for authentication
- ❑ Also, one mode of IPSec relies on IP address for authentication

TCP 3-way Handshake



- ❑ Initial sequence numbers: SEQ a and SEQ b
 - Supposed to be selected at random
- ❑ If not, might have problems...

TCP Authentication Attack



Trudy

1. SYN, SEQ = t (as Trudy)

2. SYN, ACK = t+1, SEQ = b₁

⋮

3. SYN, SEQ = t (as Alice)

5. ACK = b₂+1, data



Bob


5.
5.
5.
5.

Alice




4. SYN, ACK = t+1, SEQ = b₂

TCP Authentication Attack



Random SEQ numbers



Initial SEQ numbers
Mac OS X

- ❑ If initial SEQ numbers not very random...
- ❑ ...possible to guess initial SEQ number...
- ❑ ...and previous attack will succeed

TCP Authentication Attack

- ❑ Trudy cannot see what Bob sends, but she can send packets to Bob, while posing as **Alice**
- ❑ Trudy must prevent Alice from receiving Bob's response (or else connection will terminate)
- ❑ If **password** (or other authentication) required, this attack fails
- ❑ If TCP connection is relied on for authentication, then attack might succeed
- ❑ **Bad idea** to rely on TCP for authentication

Best Authentication Protocol?

- ❑ It depends on...
 - The sensitivity of the application/data
 - The delay that is tolerable
 - The cost (computation) that is tolerable
 - What crypto is supported (public key, symmetric key, ...)
 - Whether mutual authentication is required
 - Whether PFS, anonymity, etc., are concern
- ❑ ...and possibly other factors



Secure Shell (SSH)

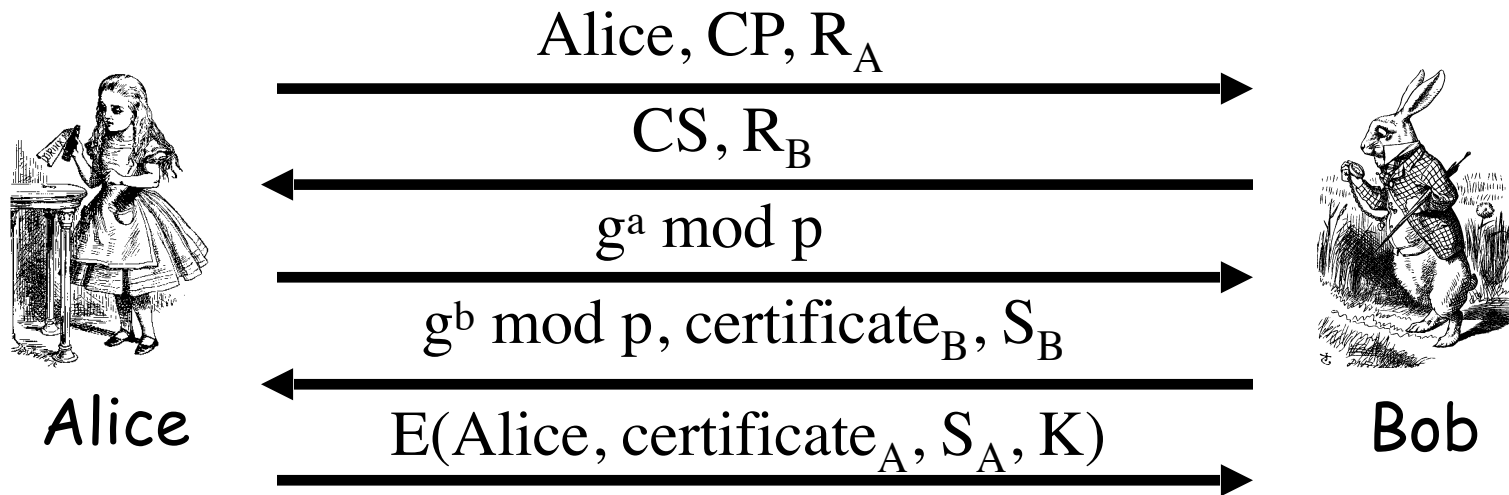
SSH

- ❑ Creates a "secure tunnel"
- ❑ Insecure command sent thru SSH "tunnel" are then secure
- ❑ SSH used with things like rlogin
 - Why is rlogin insecure without SSH?
 - Why is rlogin secure with SSH?
- ❑ SSH is a relatively simple protocol

SSH


- ❑ SSH authentication can be based on:
 - Public keys, or
 - Digital certificates, or
 - Passwords
- ❑ Here, we consider **certificate** mode
- ❑ We consider slightly simplified SSH...

Simplified SSH



- $\text{CP} = \text{"crypto proposed"}$, and $\text{CS} = \text{"crypto selected"}$
- $H = h(\text{Alice, Bob, CP, CS, } R_A, R_B, g^a \bmod p, g^b \bmod p, g^{ab} \bmod p)$
- $S_B = [H]_{\text{Bob}}$
- $S_A = [H, \text{Alice, certificate}_A]_{\text{Alice}}$
- $K = g^{ab} \bmod p$

MiM Attack on SSH?



❑ Where does this attack fail?

❑ Alice computes


$$H_a = h(\text{Alice}, \text{Bob}, \text{CP}, \text{CS}, R_A, R_B, g^a \bmod p, g^t \bmod p, g^{at} \bmod p)$$

❑ But Bob signs

$$H_b = h(\text{Alice}, \text{Bob}, \text{CP}, \text{CS}, R_A, R_B, g^t \bmod p, g^b \bmod p, g^{bt} \bmod p)$$




Secure Socket Layer



Socket layer

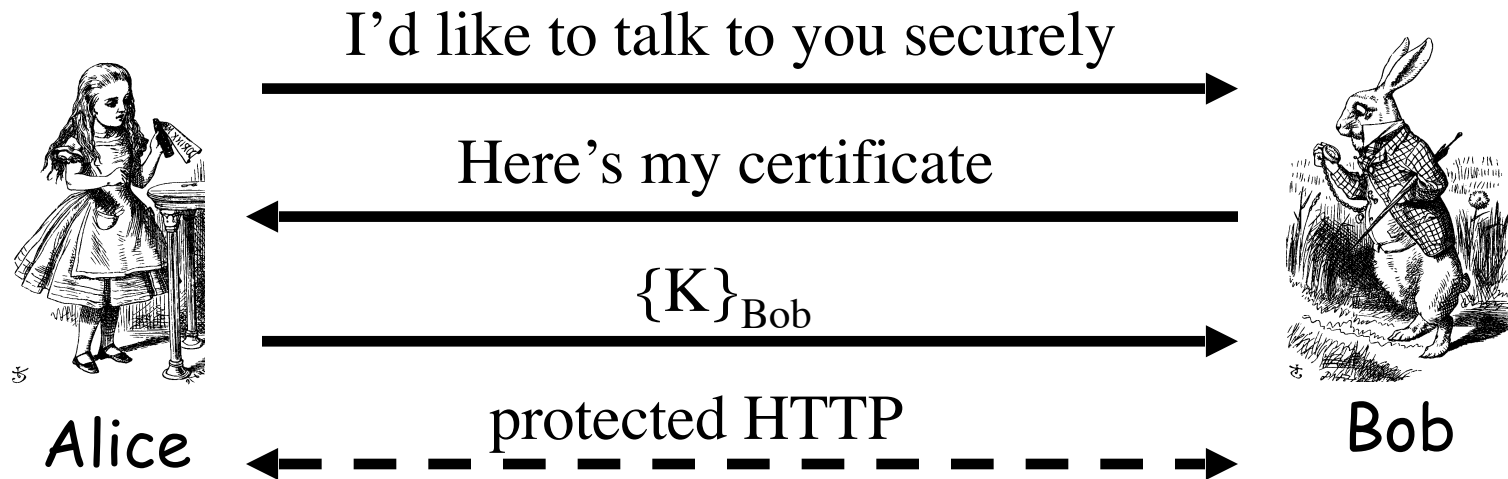
- ❑ "Socket layer" lives between application and transport layers
- ❑ SSL usually between HTTP and TCP



What is SSL?

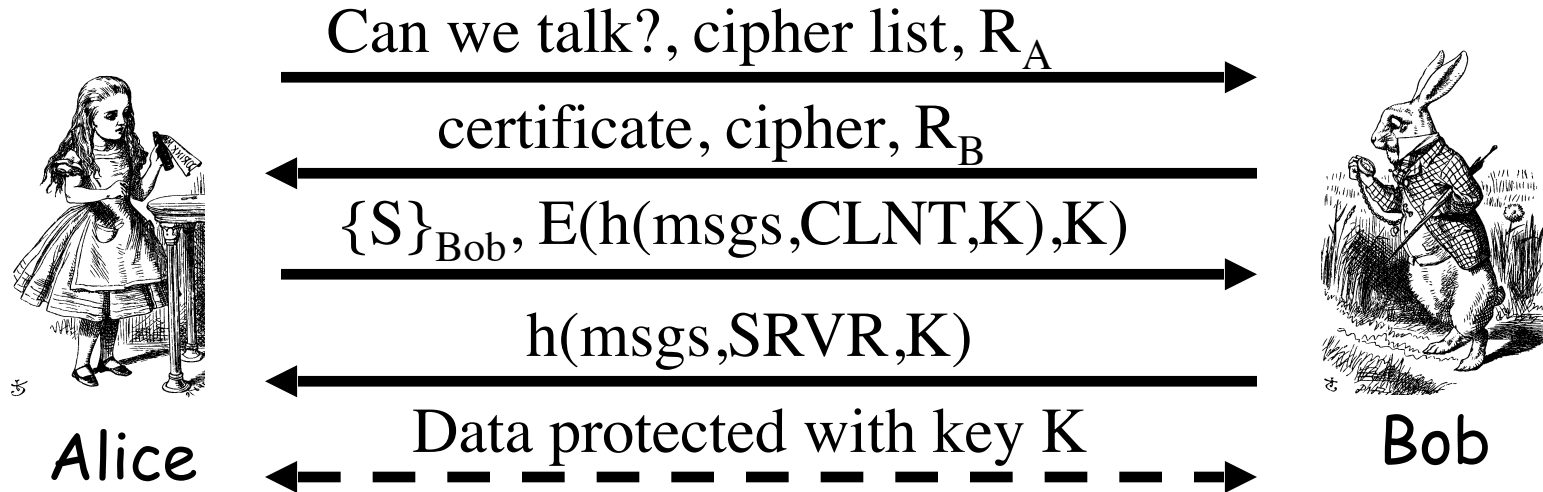
- ❑ SSL is the protocol used for majority of secure Internet transactions today
- ❑ For example, if you want to buy a book from amazon.com...
 - You want to be sure you are dealing with Amazon (**authentication**)
 - Your credit card information must be protected in transit (**confidentiality** and/or **integrity**)
 - As long as you have money, Amazon does not really care who you are...
 - ...so, no need for mutual authentication

Simple SSL-like Protocol



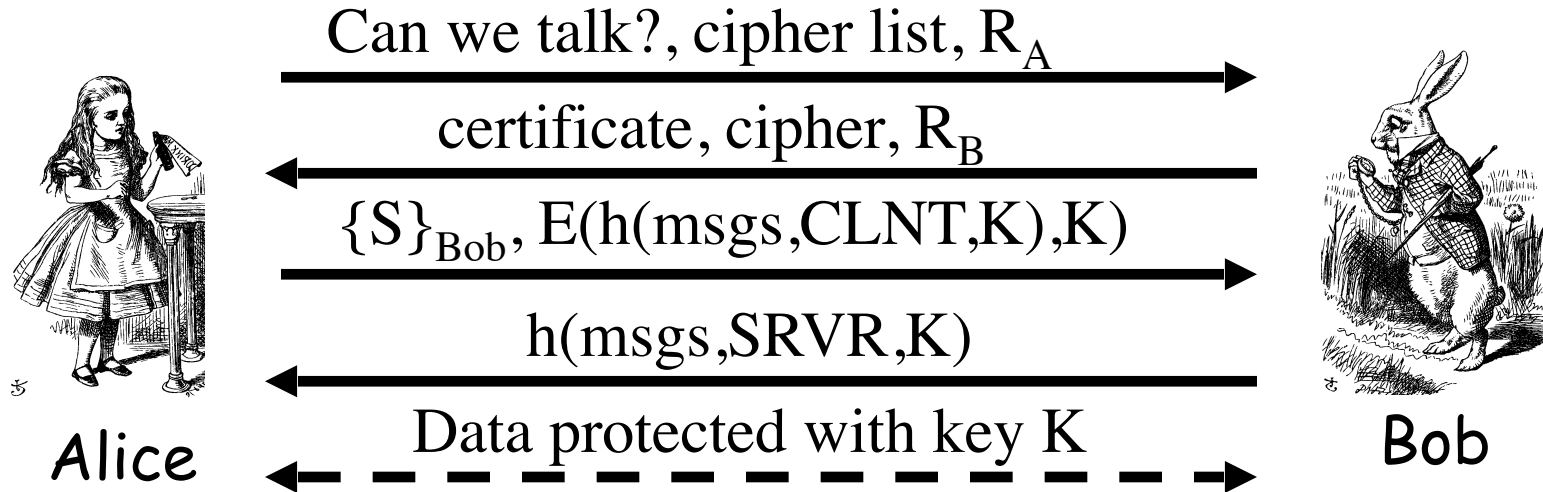
- ❑ Is Alice sure she's talking to Bob?
- ❑ Is Bob sure he's talking to Alice?

Simplified SSL Protocol



- S is the so-called **pre-master secret**
- $K = h(S, R_A, R_B)$
- "msgs" means all previous messages
- CLNT and SRVR are constants

Simplified SSL Protocol




- ❑ **Q:** Why is $h(msgs, CLNT, K)$ encrypted?
- ❑ **A:** Apparently, it adds no security...

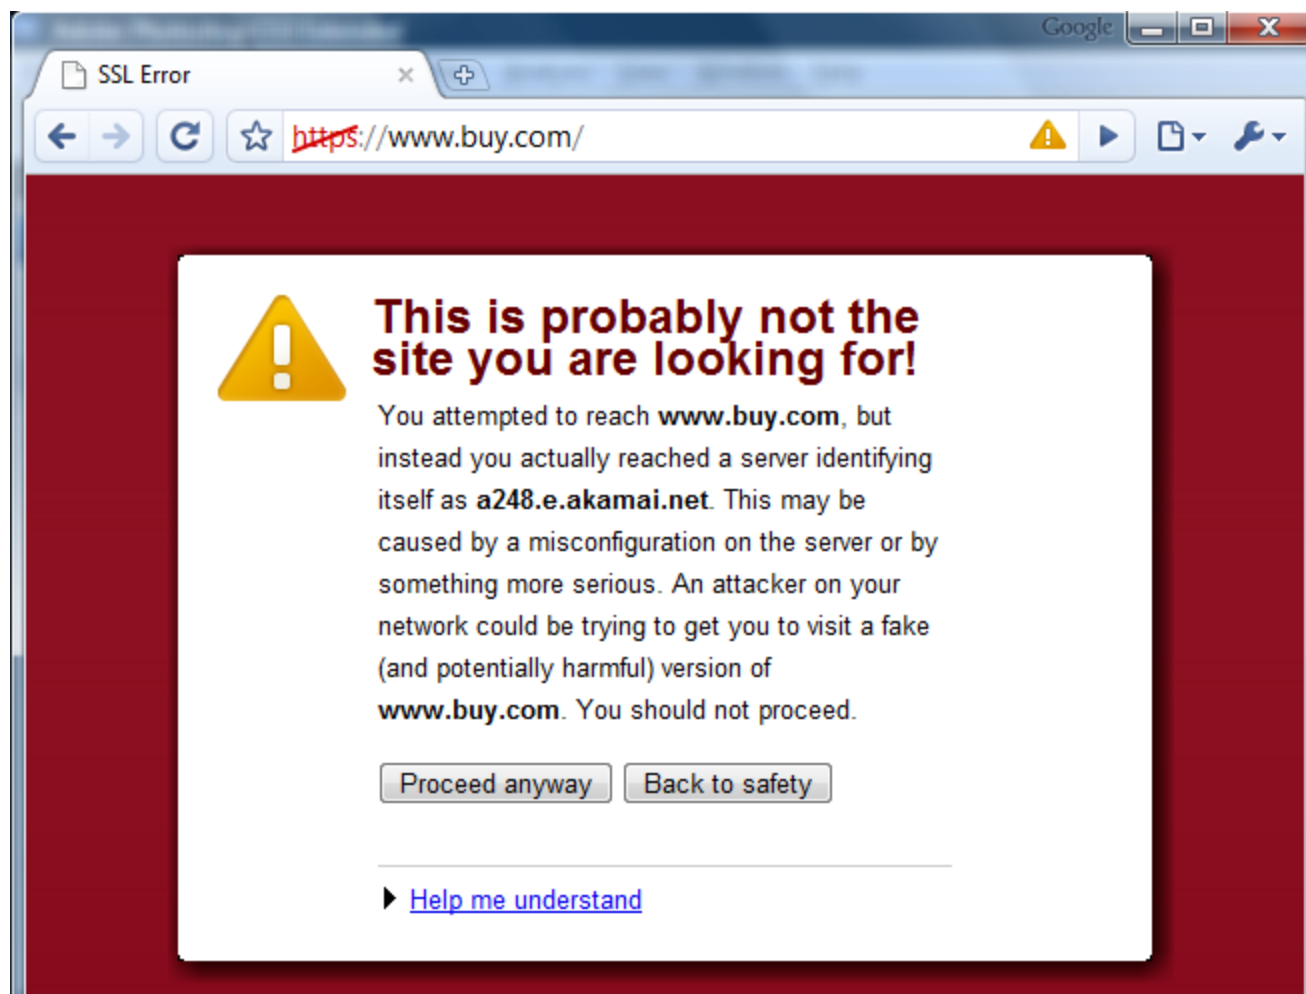
SSL Authentication

- ❑ Alice authenticates Bob, not vice-versa
 - How does client authenticate server?
 - Why would server not authenticate client?
- ❑ Mutual authentication is possible: Bob sends **certificate request** in message 2
 - Then client must have a valid certificate
 - But, if server wants to authenticate client, server could instead require password

SSL MiM Attack?



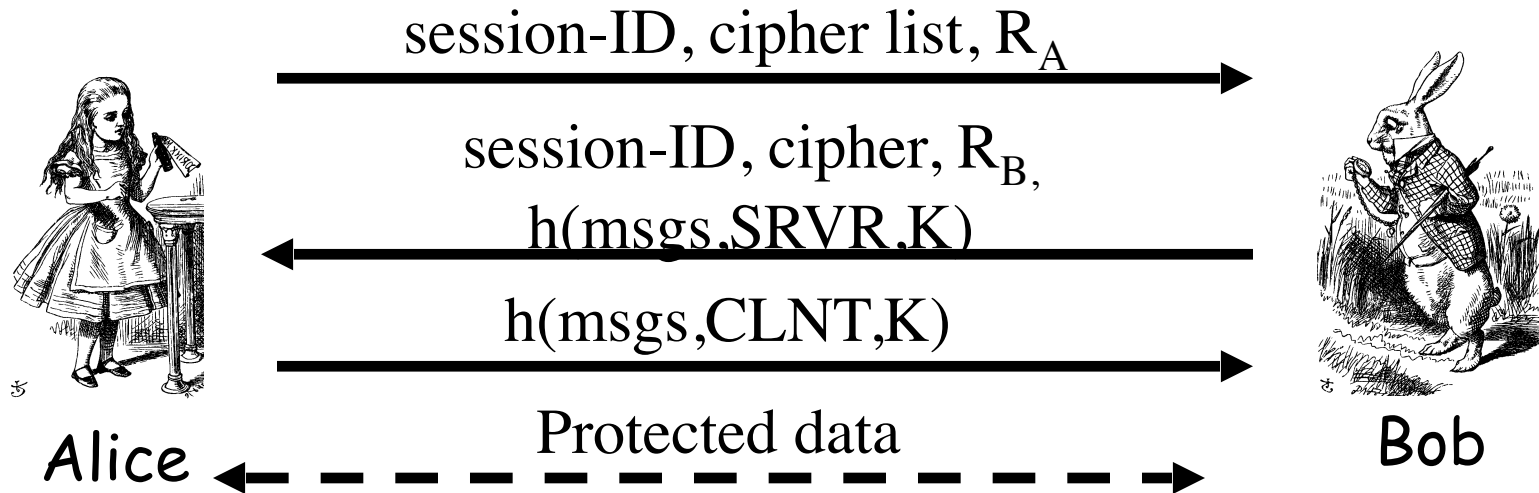
- ❑ **Q:** What prevents this MiM "attack"?
- ❑ **A:** Bob's certificate must be signed by a certificate authority (CA)
- ❑ What does browser do if signature not valid?
- ❑ What does user do when browser complains?



SSL Sessions vs Connections

- ❑ SSL **session** is established as shown on previous slides
- ❑ SSL designed for use with HTTP 1.0
- ❑ HTTP 1.0 often opens multiple simultaneous (parallel) **connections**
 - Multiple connections per session
- ❑ **SSL session is costly**, public key operations
- ❑ SSL has an efficient protocol for opening new connections **given an existing session**

SSL Connection

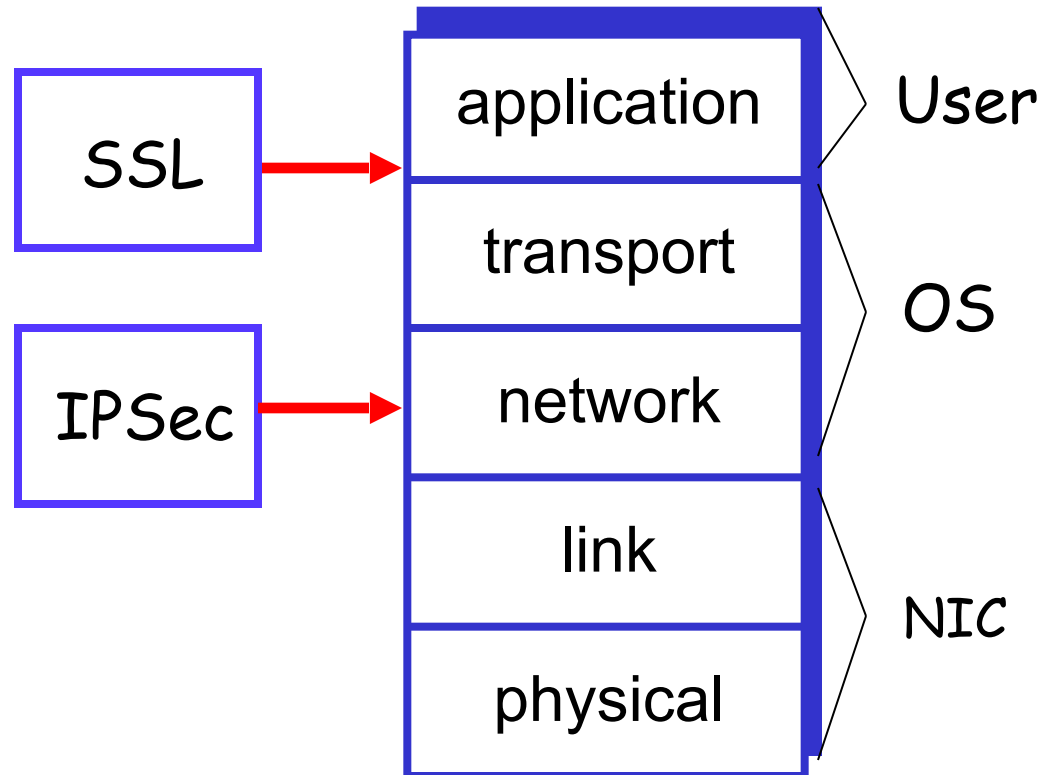


- ❑ Assuming SSL **session** exists
- ❑ So, S is already known to Alice and Bob
- ❑ Both sides must remember session-ID
- ❑ Again, $K = h(S, R_A, R_B)$
- ❑ **No public key operations!** (relies on known S)

IPSec

IPSec and SSL

- ❑ IPSec lives at the network layer
- ❑ IPSec is transparent to applications



IPSec and Complexity

- ❑ IPSec is a complex protocol
- ❑ **Over-engineered**
 - Lots of (generally useless) features
- ❑ Flawed — Some significant security issues
- ❑ Interoperability is serious challenge
 - Defeats the purpose of having a standard!
- ❑ Complex
- ❑ And, did I mention, it's complex?

IKE and ESP/AH

- ❑ Two parts to IPSec...
- ❑ **IKE**: Internet Key Exchange
 - Mutual authentication
 - Establish session key
 - Two “phases” — like SSL session/connection
- ❑ **ESP/AH**
 - **ESP**: Encapsulating Security Payload — for confidentiality and/or integrity
 - **AH**: Authentication Header — integrity only

IKE

- ❑ IKE has 2 phases
 - Phase 1 — IKE security association (SA)
 - Phase 2 — AH/ESP security association
- ❑ Phase 1 is comparable to SSL **session**
- ❑ Phase 2 is comparable to SSL **connection**
- ❑ Not an obvious need for two phases in IKE
 - In the context of IPSec, that is
- ❑ If multiple Phase 2's do not occur, then it is **more** costly to have two phases!

IKE Phase 1

- ❑ 4 different “key options”
 - Public key encryption (original version)
 - Public key encryption (improved version)
 - Public key signature
 - Symmetric key
- ❑ For each of these, 2 different “modes”
 - Main mode and aggressive mode
- ❑ **There are 8 versions of IKE Phase 1!**
- ❑ Need more evidence it's over-engineered?

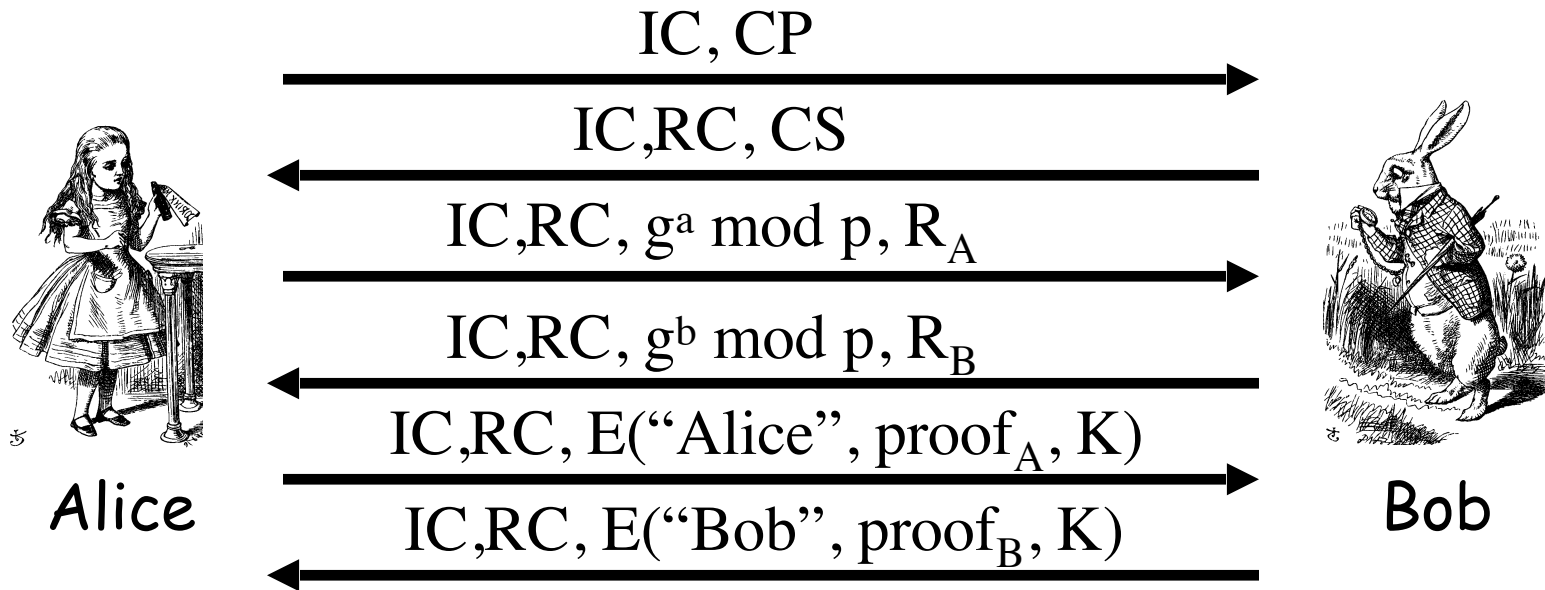
IKE Phase 1

- ❑ We discuss 6 of the 8 Phase 1 variants
 - Public key signatures (main & aggressive modes)
 - Symmetric key (main and aggressive modes)
 - Public key encryption (main and aggressive)
- ❑ Why public key encryption and public key signatures?
 - Always know your own private key
 - **May not** (initially) know other side's public key

IKE Phase 1

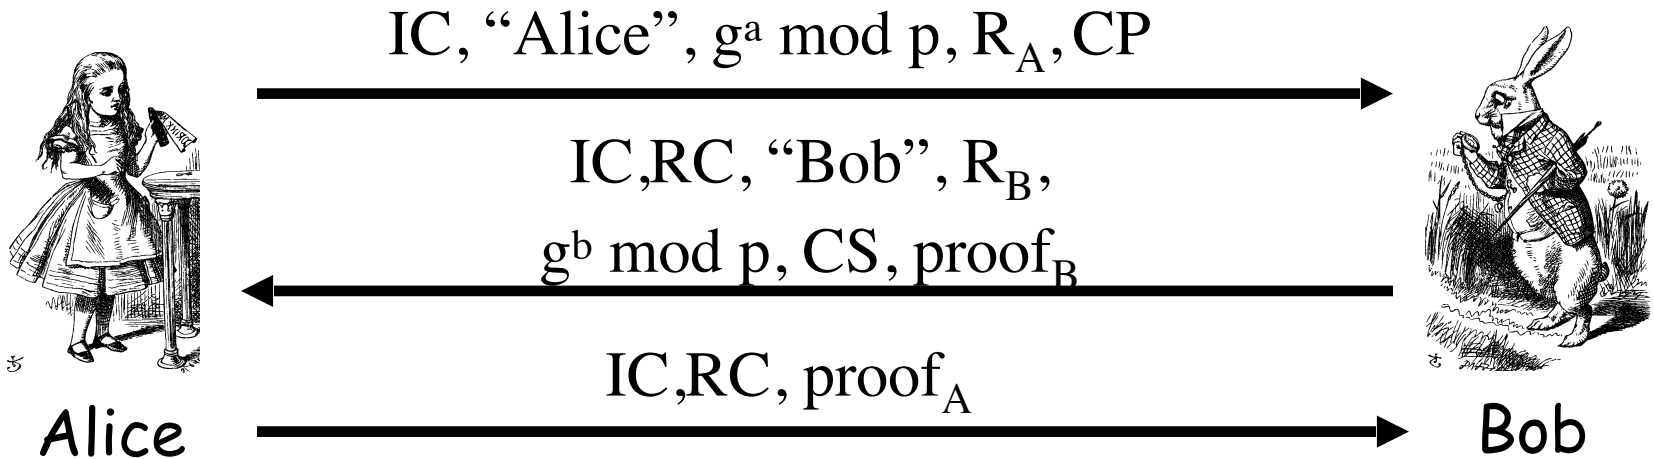
- ❑ Uses ephemeral Diffie-Hellman to establish session key
 - Provides perfect forward secrecy (PFS)
- ❑ Let a be Alice's Diffie-Hellman exponent
- ❑ Let b be Bob's Diffie-Hellman exponent
- ❑ Let g be generator and p prime
- ❑ Recall that p and g are public

IKE Phase 1: Digital Signature (Main Mode)



- CP = crypto proposed, CS = crypto selected
- IC = initiator "cookie", RC = responder "cookie"
- $K = h(IC, RC, g^{ab} \bmod p, R_A, R_B)$
- $SKEYID = h(R_A, R_B, g^{ab} \bmod p)$
- $proof_A = [h(SKEYID, g^a \bmod p, g^b \bmod p, IC, RC, CP,$
"Alice")]_{Alice}

IKE Phase 1: Public Key Signature (Aggressive Mode)




- Main differences from main mode
 - Not trying to hide identities
 - Cannot negotiate g or p

Main vs Aggressive Modes

- ❑ Main mode **MUST** be implemented
- ❑ Aggressive mode **SHOULD** be implemented
 - So, if aggressive mode is not implemented, “you should feel guilty about it”
- ❑ Might create interoperability issues
- ❑ For public key signature authentication
 - **Passive attacker** knows identities of Alice and Bob in aggressive mode, but not in main mode
 - **Active attacker** can determine Alice's and Bob's identity in main mode

IKE Phase 1: Symmetric Key (Main Mode)



□ Same as signature mode except

- K_{AB} = symmetric key shared in advance
- $K = h(IC, RC, g^{ab} \bmod p, R_A, R_B, K_{AB})$
- $SKEYID = h(K, g^{ab} \bmod p)$
- $\text{proof}_A = h(SKEYID, g^a \bmod p, g^b \bmod p, IC, RC, CP, \text{"Alice"})$

Problems with Symmetric Key (Main Mode)

❑ Catch-22

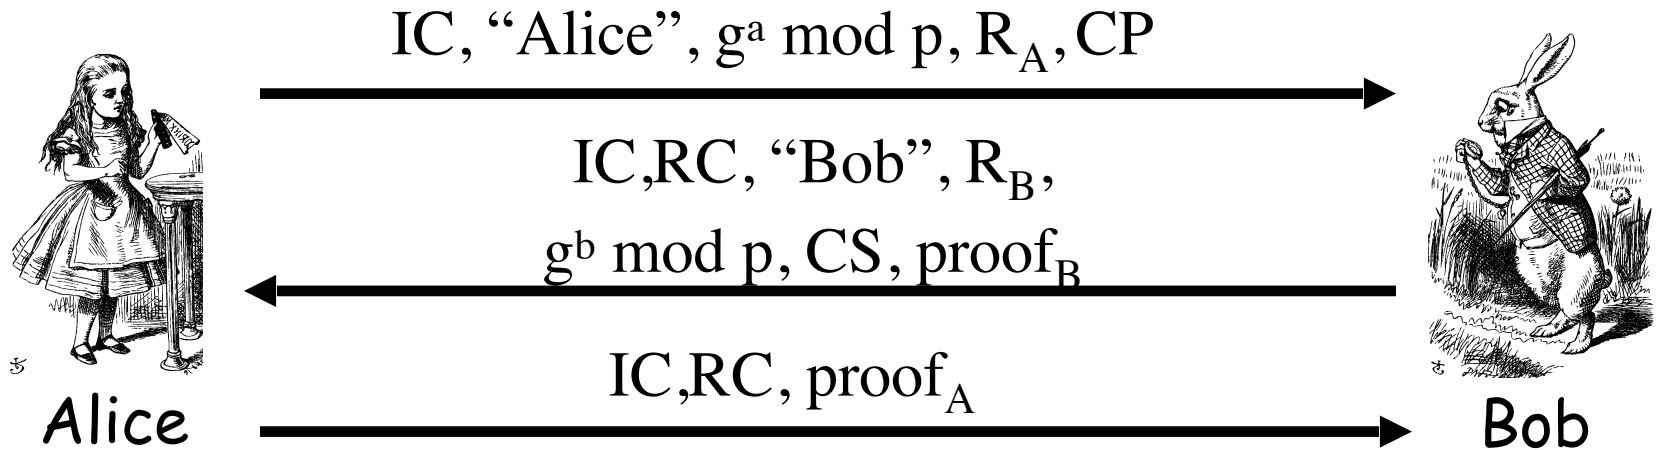
- Alice sends her ID in message 5
- Alice's ID encrypted with K
- To find K Bob must know K_{AB}
- To get K_{AB} Bob must know he's talking to Alice!

❑ Result: **Alice's IP address used as ID!**

❑ Useless mode for the "road warrior"

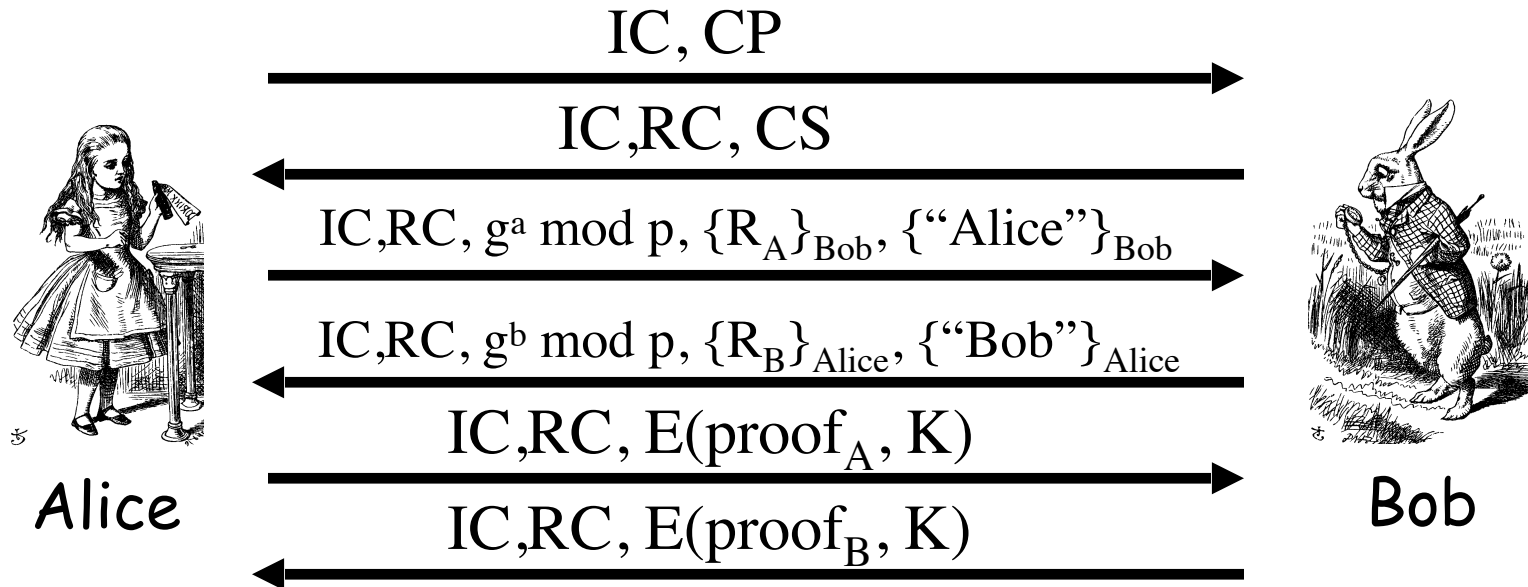
❑ Why go to all of the trouble of trying to hide identities in 6 message protocol?

IKE Phase 1: Symmetric Key (Aggressive Mode)



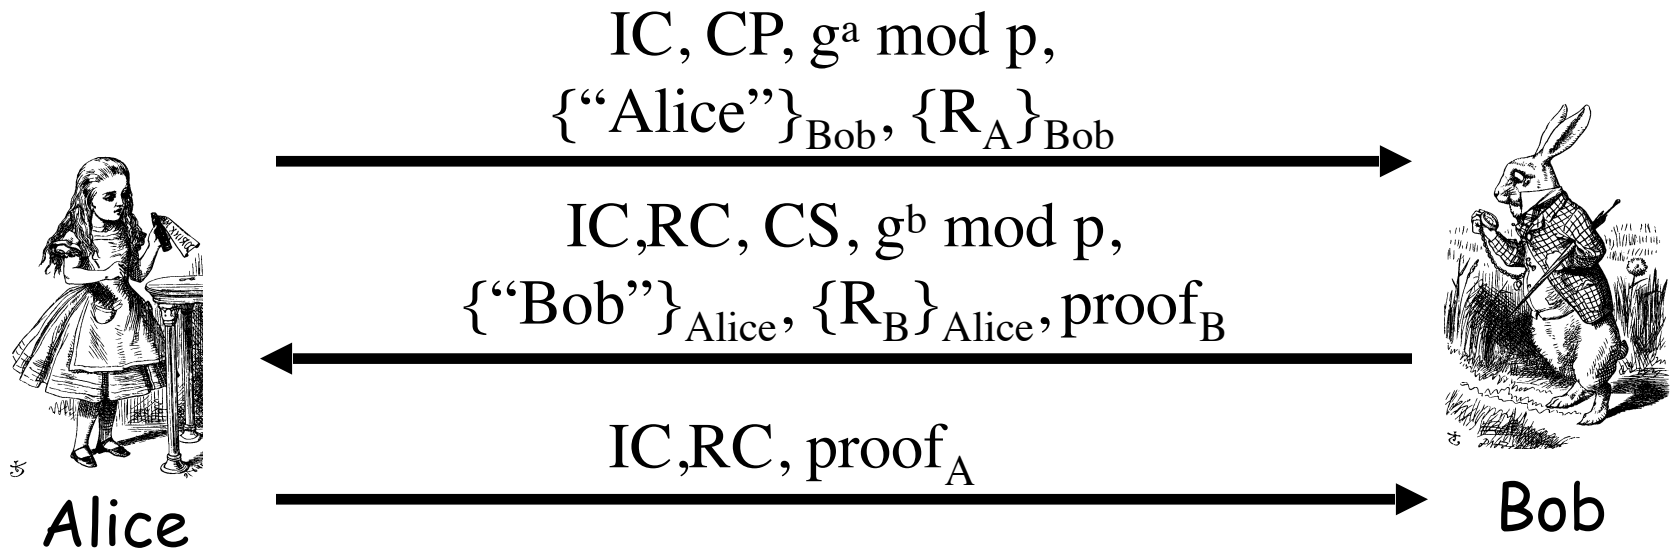
- ❑ Same format as digital signature aggressive mode
- ❑ Not trying to hide identities...
- ❑ As a result, does **not** have problems of main mode
- ❑ But does not (pretend to) hide identities

IKE Phase 1: Public Key Encryption (Main Mode)



- ❑ CP = crypto proposed, CS = crypto selected
- ❑ IC = initiator “cookie”, RC = responder “cookie”
- ❑ $K = h(IC, RC, g^{ab} \bmod p, R_A, R_B)$
- ❑ $SKEYID = h(R_A, R_B, g^{ab} \bmod p)$
- ❑ $\text{proof}_A = h(SKEYID, g^a \bmod p, g^b \bmod p, IC, RC, CP, \text{"Alice"})$

IKE Phase 1: Public Key Encryption (Aggressive Mode)

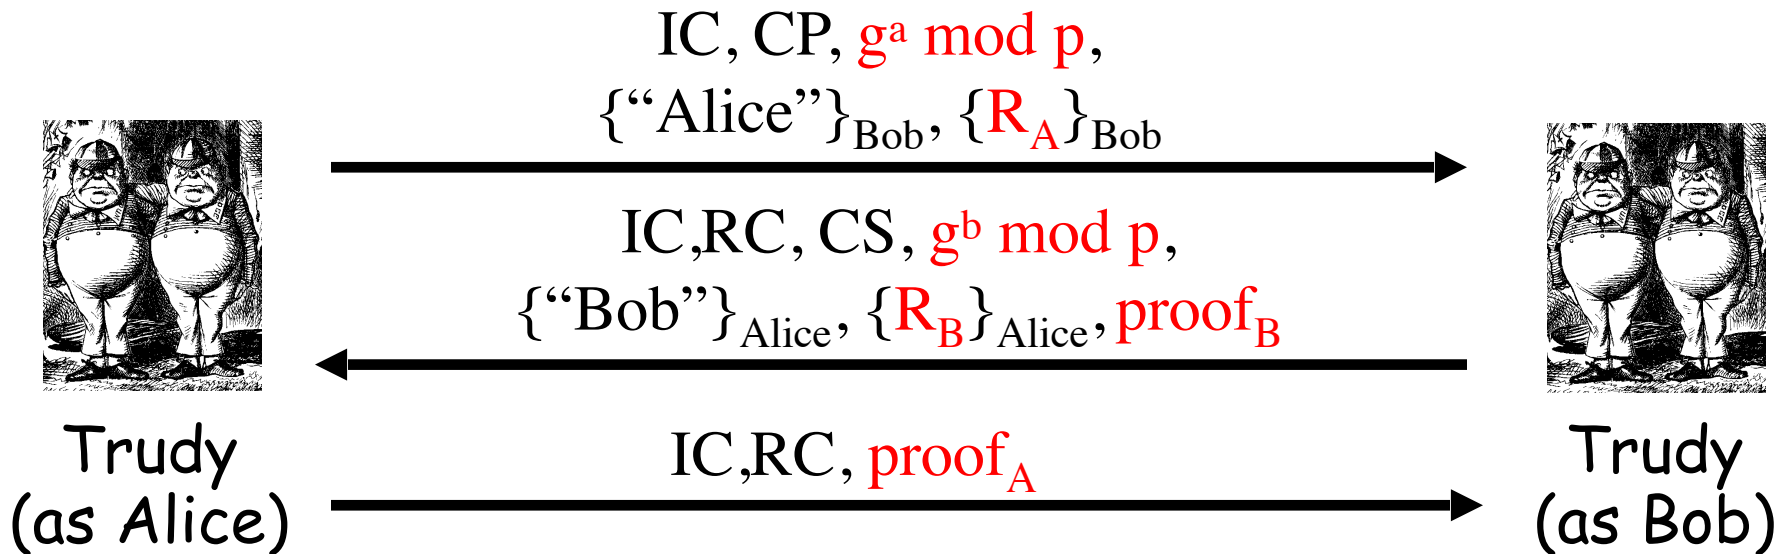


- ❑ $K, proof_A, proof_B$ computed as in main mode
- ❑ Note that identities are hidden
 - The only aggressive mode to hide identities
 - So, why have a main mode?

Public Key Encryption Issue?

- ❑ In public key encryption, aggressive mode...
- ❑ Suppose **Trudy** generates
 - Exponents **a** and **b**
 - Nonces **R_A** and **R_B**
- ❑ Trudy can compute “valid” keys and proofs:
 $g^{ab} \bmod p$, K, SKEYID, proof_A and proof_B
- ❑ All of this also works in main mode

Public Key Encryption Issue?



- ❑ Trudy can create messages that appears to be between Alice and Bob
- ❑ Appears valid to any observer, including Alice and Bob!

Plausible Deniability

- ❑ Trudy can create fake “conversation” that appears to be between Alice and Bob
 - Appears valid, even to Alice and Bob!
- ❑ A security **failure**?
- ❑ In IPSec public key option, it is a **feature**...
 - **Plausible deniability**: Alice and Bob can deny that any conversation took place!
- ❑ In some cases it might create a problem
 - E.g., if Alice makes a purchase from Bob, she could later repudiate it (unless she had signed)

IKE Phase 1 "Cookies"

- ❑ IC and RC — cookies (or "anti-clogging tokens") supposed to prevent DoS attacks
 - No relation to Web cookies
- ❑ To reduce DoS threats, Bob wants to remain **stateless** as long as possible
- ❑ But Bob must remember CP from message 1 (required for proof of identity in message 6)
- ❑ Bob must keep state from 1st message on
 - So, these "cookies" offer little DoS protection

Kerberos



Cerberos, fiera crudele e diversa,
con tre gola caninamente latra
sopra la gente che quivi è sommersa...

Kerberos

- ❑ In Greek mythology, Kerberos is 3-headed dog that guards entrance to Hades
 - “Wouldn’t it make more sense to guard the exit?”
- ❑ In security, Kerberos is an authentication protocol based on symmetric key crypto
 - Originated at MIT
 - Based on Needham and Schroeder protocol
 - Relies on a **Trusted Third Party (TTP)**

Motivation for Kerberos

- ❑ Authentication using public keys
 - N users \Rightarrow N key pairs
- ❑ Authentication using symmetric keys
 - N users requires (on the order of) N^2 keys
- ❑ Symmetric key case **does not scale**
- ❑ Kerberos based on symmetric keys but only requires N keys for N users
 - Security depends on TTP
 - + No PKI is needed

Kerberos KDC

- ❑ Kerberos **Key Distribution Center** or **KDC**
 - KDC acts as the TTP
 - TTP is trusted, so it must not be compromised
- ❑ KDC shares symmetric key K_A with Alice, key K_B with Bob, key K_C with Carol, etc.
- ❑ And a master key K_{KDC} known **only** to KDC
- ❑ KDC enables authentication, session keys
 - Session key for confidentiality and integrity
- ❑ In practice, crypto algorithm is DES


Kerberos Tickets

- ❑ KDC issue **tickets** containing info needed to access network resources
- ❑ KDC also issues **Ticket-Granting Tickets** or **TGTs** that are used to obtain tickets
- ❑ Each TGT contains
 - Session key
 - User's ID
 - Expiration time
- ❑ Every TGT is encrypted with K_{KDC}
 - So, TGT can only be read by the KDC

Kerberized Login


- ❑ Alice enters her password
- ❑ Then Alice's computer does following:
 - Derives K_A from Alice's password
 - Uses K_A to get TGT for Alice from KDC
- ❑ Alice then uses her TGT (credentials) to securely access network resources
- ❑ **Plus:** Security is transparent to Alice
- ❑ **Minus:** KDC must be secure — it's trusted!

Kerberized Login



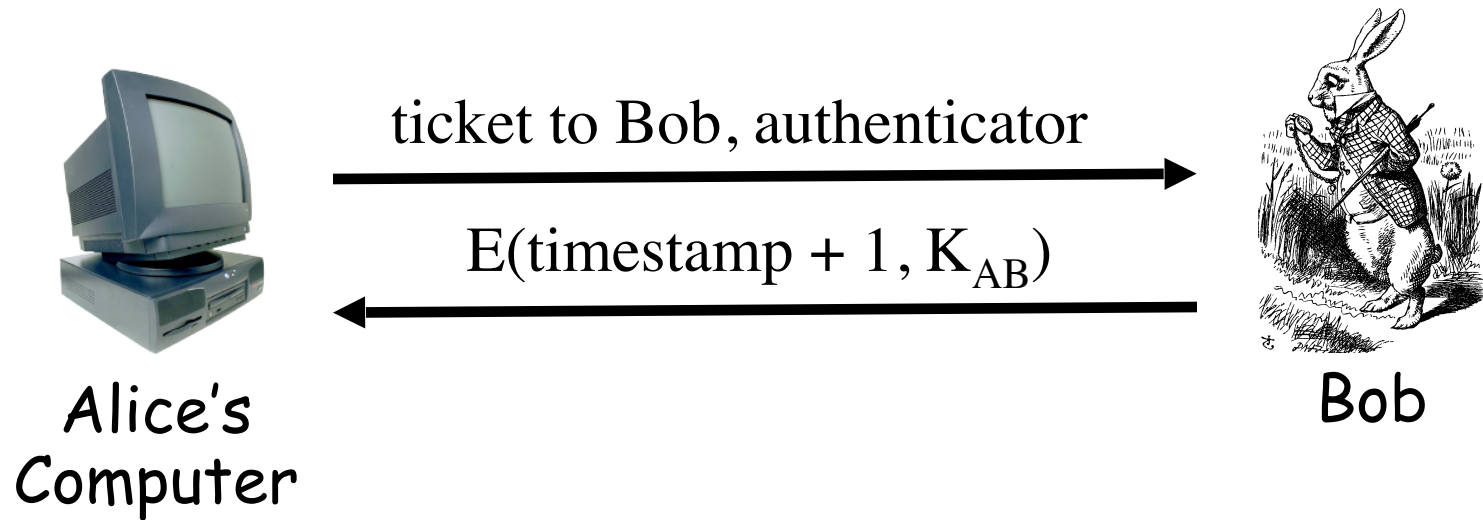
- ❑ Key $K_A = h(\text{Alice's password})$
- ❑ KDC creates **session key** S_A
- ❑ Alice's computer decrypts S_A and TGT
 - Then it forgets K_A
- ❑ $TGT = E(\text{"Alice"}, S_A, K_{KDC})$

Alice Requests "Ticket to Bob"



- ❑ REQUEST = (TGT, authenticator)
 - authenticator = $E(\text{timestamp}, S_A)$
- ❑ REPLY = $E(\text{"Bob"}, K_{AB}, \text{ticket to Bob}, S_A)$
 - ticket to Bob = $E(\text{"Alice"}, K_{AB}, K_B)$
- ❑ KDC gets S_A from TGT to verify timestamp

Alice Uses Ticket to Bob



- ❑ ticket to Bob = $E(\text{"Alice"}, K_{AB}, K_B)$
- ❑ authenticator = $E(\text{timestamp}, K_{AB})$
- ❑ Bob decrypts "ticket to Bob" to get K_{AB} which he then uses to verify timestamp

Kerberos

- ❑ Key S_A used in authentication
 - For confidentiality/integrity
- ❑ Timestamps for authentication and replay protection
- ❑ Recall, that timestamps...
 - Reduce the number of messages — like a nonce that is known in advance
 - But, “time” is a security-critical parameter

Questions about Kerberos

- When Alice logs in, KDC sends $E(S_A, TGT, K_A)$ where $TGT = E(\text{"Alice"}, S_A, K_{KDC})$

Q: Why is TGT encrypted with K_A ?

A: Enables Alice to be anonymous when she later uses her TGT to request a ticket

- In Alice's "Kerberized" login to Bob, why can Alice remain anonymous?
- Why is "ticket to Bob" sent to Alice?
 - Why doesn't KDC send it directly to Bob?

Kerberos Alternatives

- ❑ Could have Alice's computer remember password and use that for authentication?
 - Then no KDC required
 - But hard to protect passwords
 - Also, does not scale
- ❑ Could have KDC remember session key instead of putting it in a TGT?
 - Then no need for TGT
 - But **stateless** KDC is major feature of Kerberos

Kerberos Keys

- ❑ In Kerberos, $K_A = h(\text{Alice's password})$
- ❑ Could instead generate random K_A
 - Compute $K_h = h(\text{Alice's password})$
 - And Alice's computer stores $E(K_A, K_h)$
- ❑ Then K_A need not change when Alice changes her password
 - But $E(K_A, K_h)$ must be stored on computer
- ❑ This alternative approach is often used
 - But not in Kerberos