AI is a broader field (i.e.: the big umbrella) that contains several subfields such as machine learning, robotics, and computer vision



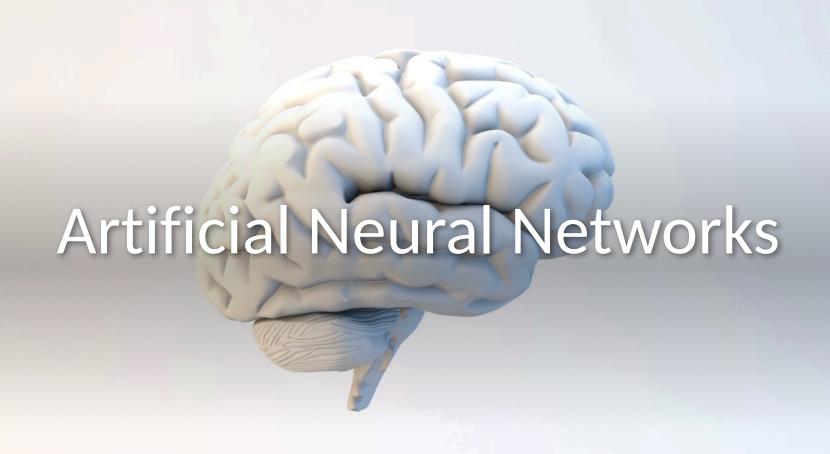
MACHINE LEARNING

DEEP LEARNING

Machine Learning is a subfield of Artificial Intelligence that enables machines to improve at a given task with experience

Deep Learning is a subfield of ML, where the algorithm tries to mathematically mimic a biological neural network.

PLUS: it "learns" the features on its own



Black box analogy



Training

There are 4 parts to understand

- Data
- Model
- ☐ Optimization Algorithm
- ☐ Objective Function

Training -Data

```
[{"index": 1, "x": 0.59857177734375, "y": 6.53887939453125, "z": 5.97772216796875}, {"index": 1, "x":
0.59857177734375, "y": 6.53887939453125, "z": 5.97772216796875}, {"index": 2, "x": 0.951507568359375,
"y": 6.8403778076171875, "z": 5.99566650390625}, {"index": 3, "x": 0.9012603759765625, "y":
6.699188232421875, "z": 5.9083251953125}, {"index": 4, "x": 0.7373504638671875, "y": 6.416839599609375,
"z": 5.44891357421875}, {"index": 5, "x": 0.34613037109375, "y": 6.3546295166015625, "z":
5.4130096435546875}, {"index": 6, "x": 0.846221923828125, "y": 6.6429595947265625, "z": 5.409423828125},
{"index": 7, "x": 0.92877197265625, "y": 6.609466552734375, "z": 5.5182952880859375}, {"index": 8, "x":
0.8593902587890625, "y": 6.4419708251953125, "z": 5.543426513671875}, {"index": 9, "x":
1.3116302490234375, "y": 6.421630859375, "z": 5.3484039306640625}, {"index": 10, "x": 1.54852294921875,
"y": 6.263702392578125, "z": 5.152191162109375}, {"index": 11, "x": 1.1931915283203125, "y":
6.146453857421875, "z": 5.0612640380859375}, {"index": 12, "x": 0.78759765625, "y": 5.9825439453125, "z"
: 5.02178955078125}, {"index": 13, "x": 0.4741363525390625, "y": 5.6798553466796875, "z":
5.2132110595703125}, {"index": 14, "x": 0.3485260009765625, "y": 5.3209228515625, "z":
5.3424224853515625}, {"index": 15, "x": 0.374847412109375, "y": 5.2993927001953125, "z":
5.4178009033203125}, {"index": 16, "x": 0.5543060302734375, "y": 5.454925537109375, "z":
5.122283935546875}, {"index": 17, "x": 0.5315704345703125, "y": 5.7791595458984375, "z":
4.77532958984375}, {"index": 18, "x": 0.623687744140625, "y": 6.0184326171875, "z": 4.4869842529296875},
{"index": 19, "x": 0.5734405517578125, "y": 6.1871337890625, "z": 4.118499755859375}, {"index": 20, "x"
: 0.1032562255859375, "y": 6.51495361328125, "z": 3.9641571044921875}, {"index": 21, "x": -
0.2425079345703125, "y": 6.9169464111328125, "z": 4.1567840576171875}, {"index": 22, "x": -
0.2544708251953125, "y": 7.29620361328125, "z": 4.68798828125}, {"index": 23, "x": 0.516021728515625,
"y": 7.6264190673828125, "z": 5.0600738525390625}, {"index": 24, "x": 1.59637451171875, "y":
7.412261962890625, "z": 4.8327484130859375}, {"index": 25, "x": 1.78302001953125, "y":
7.1334991455078125, "z": 4.7908782958984375}, {"index": 26, "x": 0.9503173828125, "y": 7.015045166015625
, "z": 5.4943695068359375}, {"index": 27, "x": 0.380828857421875, "y": 6.5472564697265625, "z":
6.4383392333984375}, {"index": 28, "x": 0.410736083984375, "y": 6.110565185546875, "z":
6.583099365234375}, {"index": 29, "x": 0.0003662109375, "y": 6.2493438720703125, "z": 6.299560546875}, {
"index": 30, "x": -0.667236328125, "y": 6.9564208984375, "z": 6.376129150390625}, {"index": 31, "x": -
```

Training - Model

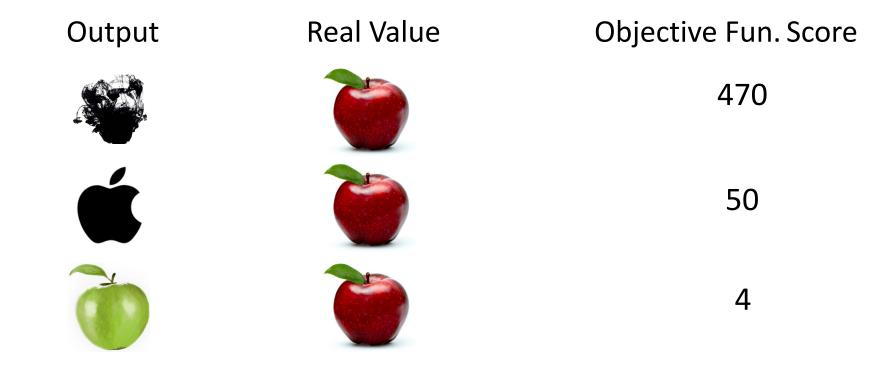
• There are several types of model.

Example: Linear

$$V_1 + V_2 + V_3 + \dots$$

Training - Objective Function

• It is the error from the labeled prediction We want to minimize it!



Training - Optimization Algorithm

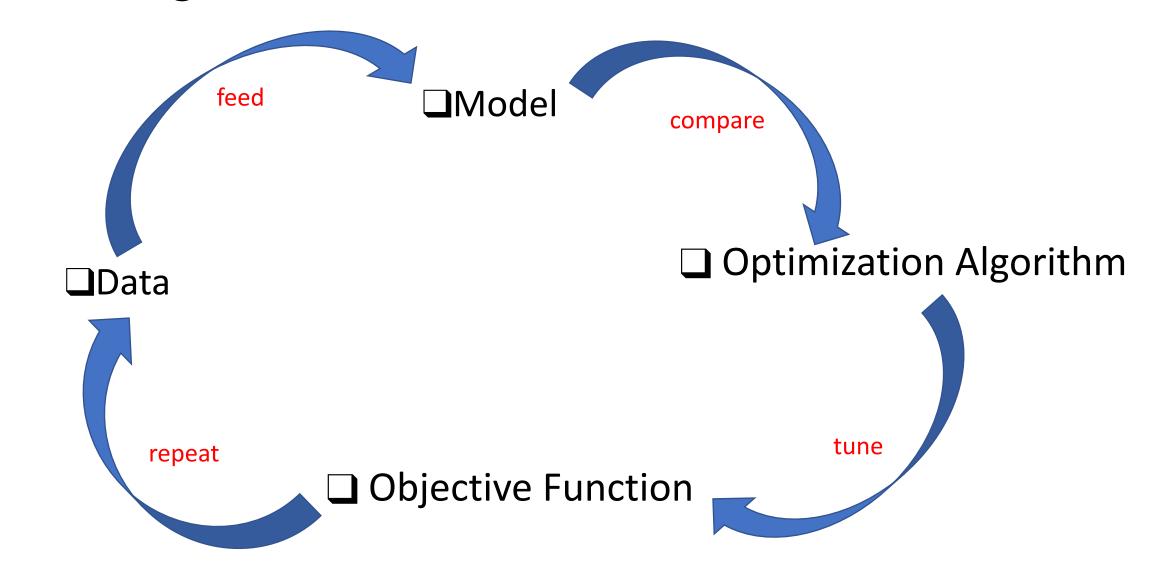
 The mechanics through which we vary the parameters of the model to optimize the Objective Function

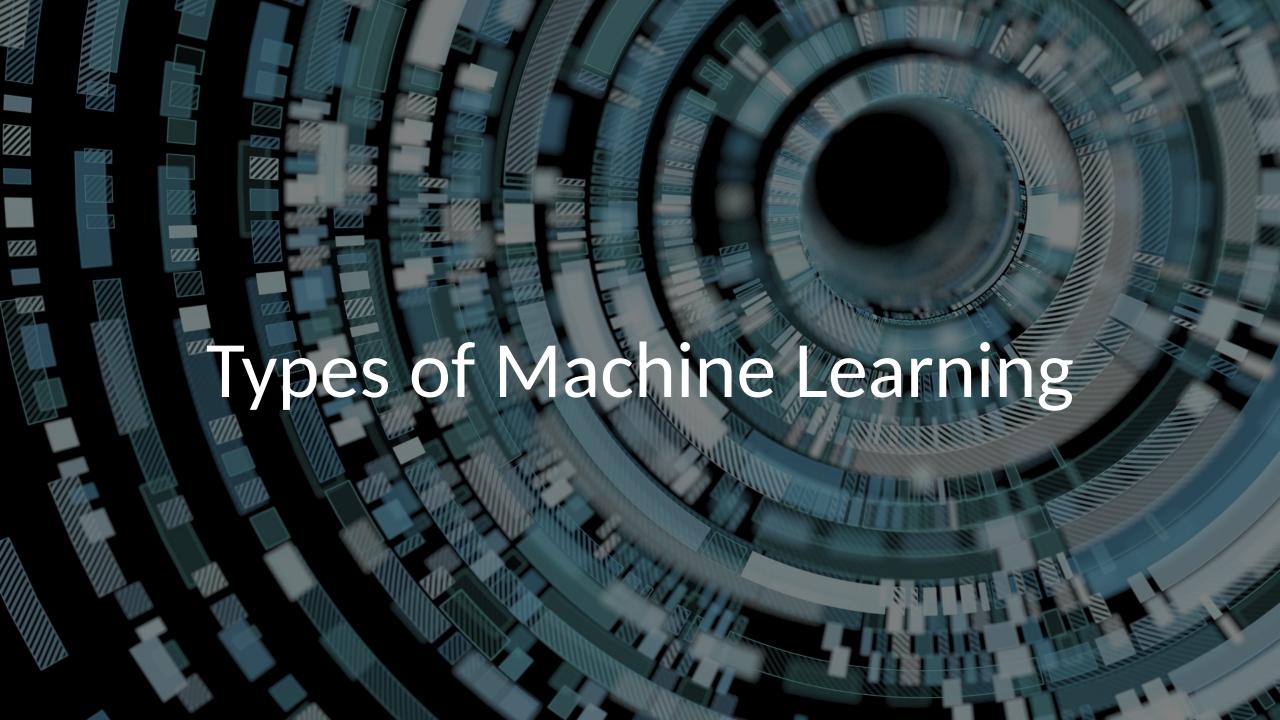
For example, we can tune the α values in a Model that looks like this:

$$\alpha_0 V_0 + \alpha_1 V_1 + \alpha_2 V_2 + \dots$$

Note that the α values can also be negative

Training





Trial-and-error

Kind of brute force attack to password guessing, but more optimized
 A trial-and-error model with feedbacks

 It could reach a "perfect" solution that is impossible to reach with a standard and schematic approach

That's one "type" of ML

Mimicking behavior

Example: Self driving vehicles

"A curve! I must turn"

"A curve, I saw others turning, I will do it too!"

That's another "type" of ML

Types of ML

□Supervised

- We have inputs, target outputs (feedback)
- The objective function in supervised learning is called loss function (also, cost or error)
- Split in **Classification**, that provides outputs (categories)
 - The labels are not ordered and cannot be compared
- And **Regression**, that provides numerical values
 - The labels can be compared

Types of ML

□Unsupervised

- Only inputs, no target outputs
- Dependence or underline logic must be found by the algorithm
- Good for data labeling (clustering)

□Reinforcement

- Sit! Good boy/girl
- The objective function is called reward function
- Good for Decision process and Reward system

Supervised

$$f(x) \rightarrow y$$

• What is *f*?

These are observations that we found:

$$f(1.7) = 10$$

$$f(0.6) = -6$$

• • •

In the classic **Linear Model**:

f(x) = x a + b

Where:

x is the **input**

a is the **coefficient** of x

b is the **intercept**

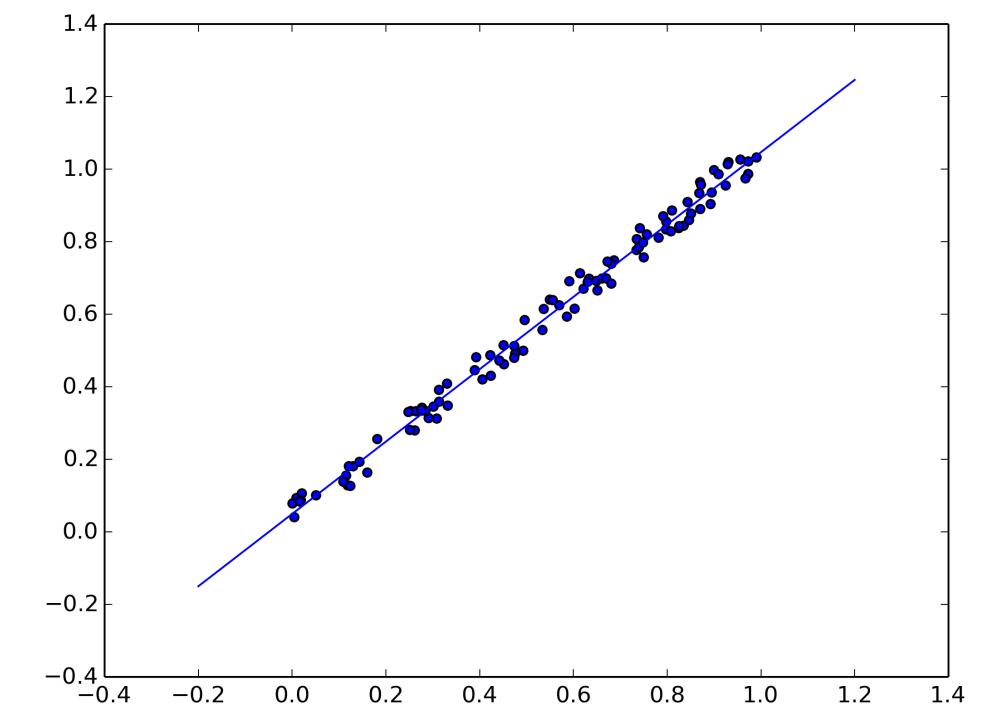
In Machine Learning, the formula is usually written as:

 $f(x) = x \, \boldsymbol{w} \, + \boldsymbol{b}$

Where:

w is the **weight**

b is the bias



f(x) = x w + b

> We want to find the values w and b so that the result f(x) is as close as possible to the observations

 Data that can be classified using a linear model is called linearly separable

f(x) = x w + b

Example:

- Imagine that you want to identify the price of apartments around SJSU
- You think that the size of the apartments is a crucial factor
- So, the value x will represent the size of such apartments

The formula may become:

f(x) = x * 100.08 + 9.98

Example:

- Imagine that you want to identify the price apartments around SJSU
- You think that the size of the apartment is a crucial factor
- So, the value x will represent the size of such apartments

Ok, but what if you think that the size is not enough to define the price? What about how close apartments are to SJSU?

- In this case we have two variables
 - O But we still use this formula:

$$f(x) = xw + b$$

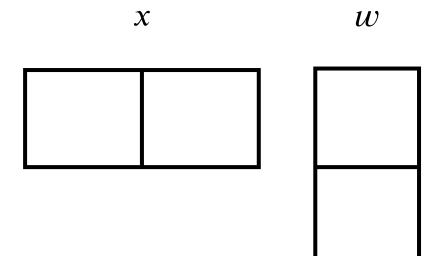
 $f(x) = x \, \boldsymbol{w} \, + \boldsymbol{b}$

ullet In fact, $oldsymbol{w}$ and $oldsymbol{x}$ are actually **vectors**

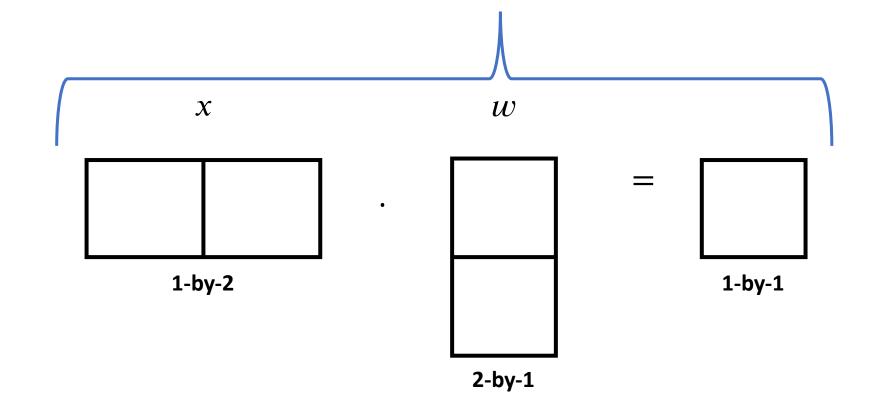
x is **1-by-2 vector**

w is **2-by-1 vector**

b is **1-by-1 vector**



 $f(x) = x \, \boldsymbol{w} \, + \boldsymbol{b}$



 $f(x) = x_1 * w_1 + x_2 * w_2 + b$

 $f(x) = \begin{bmatrix} x_1 & x_2 & & w \\ & & & \\ & & \\ & & & \\ & & & \\ & & & \\ & & & \\ & & & \\ & & & \\ & & & \\ & &$

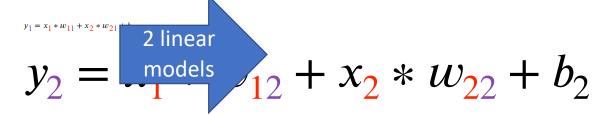
 $f(x) = x_1 * w_1 + x_2 * w_2 + b$

ullet In a realistic scenario, w_2 would be a negative weight

In fact, we expect the distance to SJSU to have a negative impact when we raise it

Multiple Inputs and Multiple Outputs

Now, you want to check the price for renting AND buying an apartment near SJSU



The number of weights is equal to:

weights = # inputs * # outputs

The number of biases is equal to:

biases = # outputs

Multiple Inputs and Multiple Outputs

$$y_1 = x_1 * w_{11} + x_2 * w_{21} + b_1$$

 $y_2 = x_1 * w_{12} + x_2 * w_{22} + b_2$

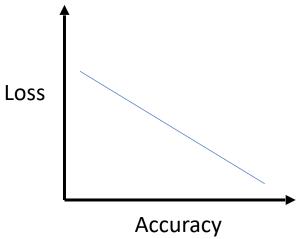
y_1	y_2	=	x_1	x_2	•	w_{11}	w_{12}	+	b_1	b_2
						w_{21}	w_{22}			

Objective Function

Objective Function

There are two categories of Objective Function: Loss and Reward

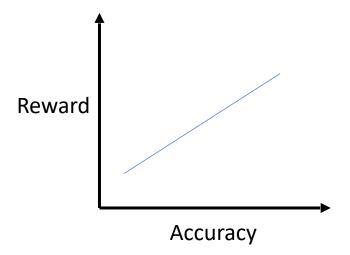
- ☐ Loss (or Cost) functions
 - The lower is the **Loss** function, the higher is the level of **accuracy**
 - Usually used in Supervised learning
 - It's nothing more than a function that has higher values for worse results and vice versa



Objective Function

There are two categories of Objective Function: Loss and Reward

- ☐ **Reward** functions
 - Usually used in Reinforcement learning
 - The exact opposite of a Loss function



Loss functions

Loss functions are used in Supervised Learning

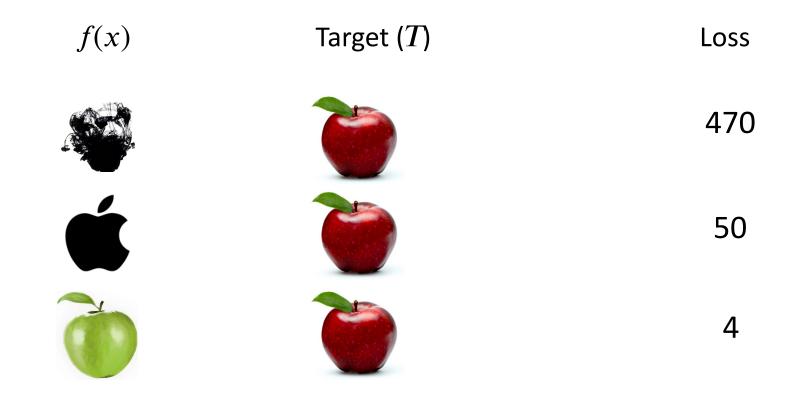
 There are two types of Supervised Learning: Regression and Classification

Two very common **Loss functions** are:

- >L2-norm (aka square loss or Least sum of squares)
 - Used for Regression
- **>**Cross-entropy
 - Used for Classification

• We need to introduce the concept of **Target** (T) Ideally, we want: y=T

The targets are the labels used to train the model



L2-norm Loss

L2-norm =
$$\sum_{i} (y_i - t_i)^2$$

- Vector norm (aka Euclidian distance) of the outputs and the targets
- The lower the error, the lower the Loss
- Note that the results of this function would always be a number

Cross-entropy Loss

Cross-entropy =
$$L(y, t) = -\sum_{i} t_{i} ln(y_{i})$$

Let's see an example

Imagine we have three categories: Apple, Pears and Pineapples

$$[0, 1, 0]$$
 $[0, 0, 1]$

Cross-entropy =
$$L(y, t) = -\sum_{i} t_{i} ln(y_{i})$$



$$Y = [0.4, 0.4, 0.2]$$

$$T = \begin{bmatrix} 0 & 1 & 0 \end{bmatrix}$$

Output of the model



$$Y = [0.1, 0.2, 0.7]$$

$$T = \begin{bmatrix} 0, 0, 1 \end{bmatrix}$$

Cross-entropy =
$$L(y, t) = -\sum t_i ln(y_i)$$



$$Y = [0.4, 0.4, 0.2]$$

$$T = \begin{bmatrix} 0, 1, 0 \end{bmatrix}$$

40% it s an Apple 40% it s a Pear 20% it s a Pineapple



$$Y = [0.1, 0.2, 0.7]$$

$$T = \begin{bmatrix} 0, 0, 1 \end{bmatrix}$$

10% it's an Apple20% it's a Pear70% it's a Pineapple

Cross-entropy =
$$L(y, t) = -\sum t_i ln(y_i)$$



$$Y = [0.4, 0.4, 0.2]$$

$$Y = [0.4, 0.4, 0.2]$$
 $0.4 \cdot \ln(0.4) - 0 \cdot \ln(0.4) - 0 \cdot \ln(0.2) = 0.92$

$$T = \begin{bmatrix} 0 , 1 , 0 \end{bmatrix}$$



$$Y = [0.1, 0.2, 0.7]$$

$$T = \begin{bmatrix} 0, 0, 1 \end{bmatrix}$$

$$Y = [0.1, 0.2, 0.7]$$
 $0.2 + \ln(0.1) - 0 + \ln(0.2) - 1 + \ln(0.7) = 0.36$

Cross-entropy =
$$L(y, t) = -\sum t_i ln(y_i)$$






$$L(y, t^{i}) = 0.92$$

$$T = \begin{bmatrix} 0, 1, 0 \end{bmatrix}$$



$$Y = [0.1, 0.2, 0.7]$$

$$T = \begin{bmatrix} 0, 0, 1 \end{bmatrix}$$



That's the best!

Optimization Algorithm

Gradient Descent

The last piece of the puzzle is the Optimization Algorithm

• The easiest Optimization Algorithm is the Gradient Descent

 Gradient Descent is a form of trial-and-error in searching for a local minimum

Computers like a lot trial-and-error approaches

 Using a "well-designed" update rule, each trial is better than the previous one

It reaches the minimum faster, without oscillation

Gradient Descent

Let's see a simple example:

$$f(x) = 5x^2 + 3x - 4$$

We want to find the MIN of the function f(x) using **Gradient Descent**

Step 1: Derivative of the function

$$f'(x) = 10x + 3$$

Step 2: Choose an arbitrary number x_0

$$x_0 = 4$$

Step 1: Derivative of the function

$$f'(x) = 10x + 3$$

Step 2: Choose an arbitrary number x_0

$$x_0 = 4$$

Step 3: Choose a value x_1 following this update rule:

 $x_{i+1} = x_i - \eta \, f'(x_i)$

$$x_1 = x_0 - \eta[10 * x_0 + 3]$$

$$x_1 = 4 - \eta [10 * 4 + 4] = 4 - 43\eta$$

Step 1: Derivative of the function

$$f'(x) = 10x + 3$$

Step 2: Choose an arbitrary number x_0

$$x_0 = 4$$

Step 3: Choose a value x_1 following this rule:

 $x_1 = 4 - 43\eta$

 $\gg \eta$ is the **learning rate** at which the ML algorithm forgets all beliefs for new ones

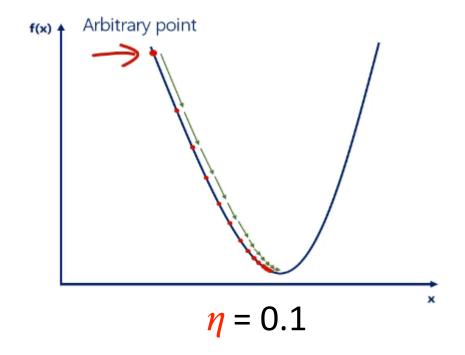
• Step 4: we keep searching for values (x_2, x_3, \ldots) until we reach a point were the values stop updating

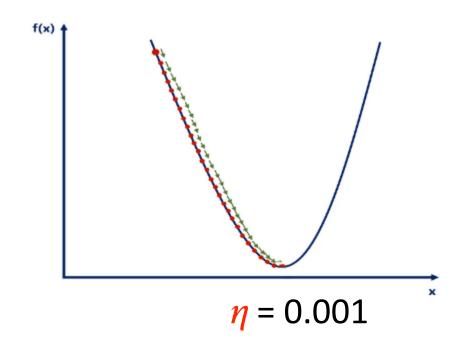
We basically reached the minimum of the function, that is, where the derivative is equal to 0

• If the derivative is 0, then the update rule: $x_{i+1} = x_i - \eta f'(x_i)$ becomes:

$$x_{i+1} = x_i - 0$$

That is, it stops updating

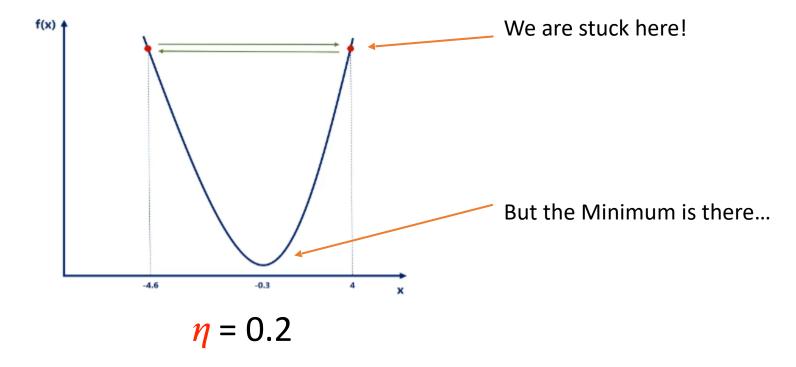




• The speed of minimization depends on η

Oscillation

- For some η values, the result of the update rule could oscillate eternally between two (or more) values.
- ➤ This is called **Oscillation**



Which η should we chose?

To choose the right η value, follow these two simple rules:

- It should be high enough, so that the closest minimum is reached in a rational amount of time
- It should be low enough, so that we don't oscillate in aeternum around the minimum

➤In other words, experiment with reasonable values...

How do I know when to stop?

 What if the results of the update rule are smaller and smaller and seem to continue forever?

• A rule of thumb is to stop when:

$$x_{i+1} - x_i = 0.001$$

This can be implemented easily adding a "breaking condition" in the loop

N-dimensional Gradient Descent

Up to now, we have:

• Data:

We select some data

Model:

$$xw + b = y$$
 [Linear Model]

• Optimization Function:

$$L(y,t) = \frac{L2 - norm}{2} = \frac{\sum_{i} (y_i - t_i)^2}{2}$$

We'll see why over 2 soon

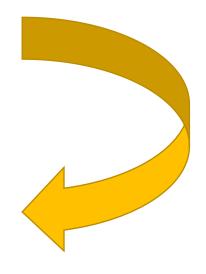
N-dimensional Gradient Descent

• The update rule:

$$x_{i+1} = x_i - \eta f'(x_i)$$

• Becomes:

$$\begin{aligned} \boldsymbol{w}_{i+1} &= \boldsymbol{w}_i - \eta \nabla_{\boldsymbol{w}} L(\boldsymbol{y}, t) \\ \boldsymbol{b}_{i+1} &= \boldsymbol{b}_i - \eta \nabla_{\boldsymbol{b}} L(\boldsymbol{y}, t) \end{aligned}$$



 $m{V}_{w}m{L}(m{y},m{t})$ is the gradient of the Loss function with respect to w_i for the weights $m{V}_{b}m{L}(m{y},m{t})$ is the gradient of the Loss function with respect to b_i for the biases

N-dimensional Gradient Descent

We update rule for N-dimensions:

$$\begin{aligned} \boldsymbol{w}_{i+1} &= \boldsymbol{w}_i - \eta \nabla_{\boldsymbol{w}} L(\boldsymbol{y}, t) \\ \boldsymbol{b}_{i+1} &= \boldsymbol{b}_i - \eta \nabla_{\boldsymbol{b}} L(\boldsymbol{y}, t) \end{aligned}$$

ullet To minimize the Loss function, we need to optimize regarding $oldsymbol{w}$ and $oldsymbol{b}$ It's still a game of tuning the weights and the biases

Optimization

• We update rule for N-dimensions:

$$\begin{aligned} \boldsymbol{w}_{i+1} &= \boldsymbol{w}_i - \eta \nabla_{\boldsymbol{w}} L(\boldsymbol{y}, \boldsymbol{t}) \\ \boldsymbol{b}_{i+1} &= \boldsymbol{b}_i - \eta \nabla_{\boldsymbol{b}} L(\boldsymbol{y}, \boldsymbol{t}) \end{aligned}$$

$$\nabla_{w}L(y,t) = \sum_{i} \nabla_{w} \frac{1}{2} (y_{i} - t_{i})^{2} = \sum_{i} x_{i} (y_{i} - t_{i}) = \sum_{i} x_{i} \delta_{i}$$

Recall that:
$$L(y,t) = \frac{\sum_{i} (y_i - t_i)^2}{2}$$
 Where w and x are matrices

 δ is the default math symbol used to measure differences

$$\nabla_{\mathbf{w}} L = \nabla_{\mathbf{w}} \frac{1}{2} \sum_{i} (y_{i} - t_{i})^{2} =$$

$$= \nabla_{\mathbf{w}} \frac{1}{2} \sum_{i} ((x_{i} w + b) - t_{i})^{2} =$$

$$= \sum_{i} \nabla_{\mathbf{w}} \frac{1}{2} (x_{i} w + b - t_{i})^{2} =$$

$$= \sum_{i} x_{i} (x_{i} w + b - t_{i}) =$$

$$= \sum_{i} x_{i} (y_{i} - t_{i}) \equiv$$

$$\equiv \sum_{i} x_{i} \delta_{i}$$

Optimization

• The update rule for N-dimensions:

$$\begin{aligned} \boldsymbol{w}_{i+1} &= \boldsymbol{w}_i - \eta \nabla_{\boldsymbol{w}} L(\boldsymbol{y}, t) \\ \boldsymbol{b}_{i+1} &= \boldsymbol{b}_i - \eta \nabla_{\boldsymbol{b}} L(\boldsymbol{y}, t) \end{aligned}$$

$$\nabla_w L(y,t) = \sum_i x_i \delta_i$$

$$\nabla_b L(y,t) = \sum_i \delta_i$$

Finally, we have:

• Data:

We select some data

• Model:

$$xw + b = y$$
 [Linear Model]

• Optimization Function:

$$L(y,t) = \frac{L2 - norm}{2} = \frac{\sum_{i} (y_i - t_i)^2}{2}$$

• Optimization Algorithm: $w_i - \eta \nabla_w L(y,t)$ $b_{i+1} = b_i - \eta \nabla_b L(y,t)$