

Setting up the working environment

Setting up the working environment

We will need to install **Anaconda**

- **Anaconda** is a software package that contains **Python** and **Jupyter Notebook**
- It also contains many useful libraries (NumPy, pandas, ...)

We will need to also install the required packages to use **TensorFlow**

Why Python?

You have a lot of choices, but usually you see projects made in **Python** or **R**. But why are we going to use Python?

Python is:

- easy to learn
- general-purpose
- very high-level
- very nice and useful packages for ML (for example, scikit learn)
- also free!

Why Python?

Isn't Python very slow?

- Yes, it is $\sim 100x$ slower than C
- But several packages are actually written in C/C++
well, even Python is...
- Also, integration with **heterogenous programming** makes it a good choice for ML

Why Jupyter Notebook?



Jupyter is a **server-client application** that:

- runs your code in your web-browser
- it's used at Google, Microsoft, IBM, ...
- incorporates several languages and allows easy sharing

Installing Anaconda

1. Go to: <https://www.anaconda.com/distribution/>
2. Select your OS
3. Download Python 3.8 (64-bit)

Anaconda Installers

Windows 	MacOS 	Linux 
Python 3.8 64-Bit Graphical Installer (466 MB) 32-Bit Graphical Installer (397 MB)	Python 3.8 64-Bit Graphical Installer (462 MB) 64-Bit Command Line Installer (454 MB)	Python 3.8 64-Bit (x86) Installer (550 MB) 64-Bit (Power8 and Power9) Installer (290 MB)

Installing Anaconda

4. For compatibility issues, check: *“Register Anaconda as my default Python 3.8”*
 - It could raise an alert if you already have Python installed
5. Don't select the Anaconda Cloud

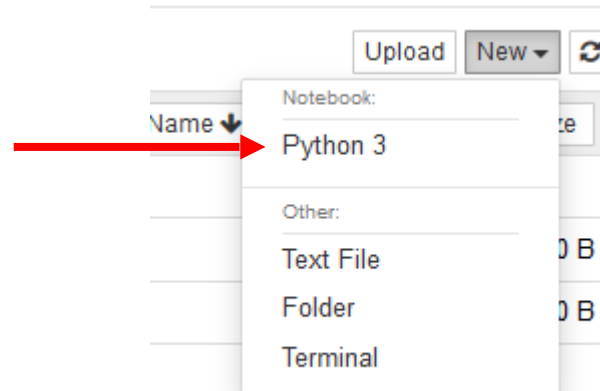
Let's start

- Run **Jupyter Notebook** (now installed with Python)

It will open in the Browser

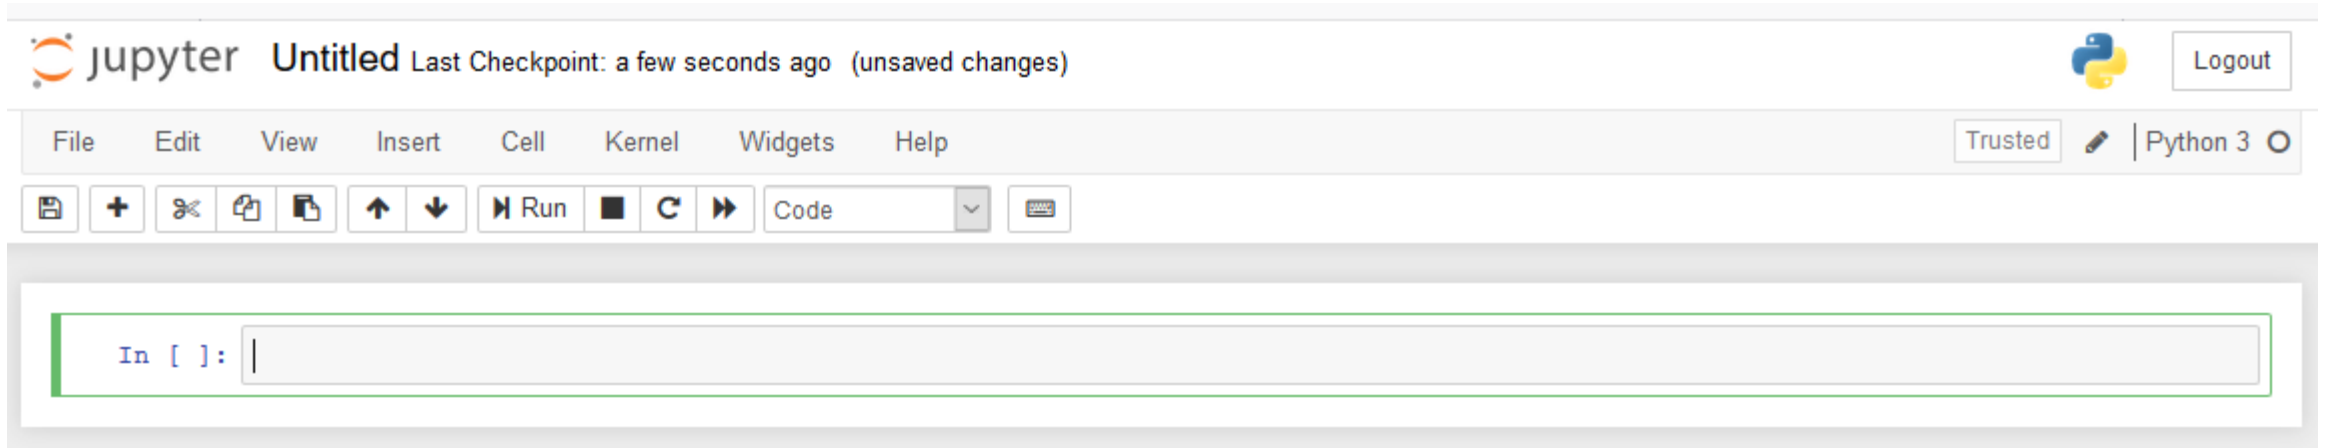
Don't close the “shell” window that might open


- Jupyter uses IPython Notebook Format (.ipynb)
- Create a new Python 3 Notebook
















Let's start

- It should open a new tab like this:



jupyter Untitled1 Last Checkpoint: a minute ago (unsaved changes)  Logout

File Edit View Insert Cell Kernel Widgets Help Trusted Python 3 

       Run    Code  

This code is not particularly useful

```
In [1]: x = 3 * 4
        x
Out[1]: 12
```

```
In [2]: y = x**2
        y
Out[2]: 144
```

```
In [ ]:
```

- It splits the code in cells that you can run individually
Very nice for debugging!

Installing TensorFlow 2.1.0

- Run “Anaconda Prompt” from your OS Menu
- Commands:

```
conda info --envs
```

```
(base) C:\Users\fabio>conda info --envs
# conda environments:
#
base                  *  C:\Users\fabio\Anaconda3
```

```
conda create --name env_name python=3
```

```
(base) C:\Users\fabio>conda create --name Python3-TensorFlow2 python=3
```

The following NEW packages will be INSTALLED:

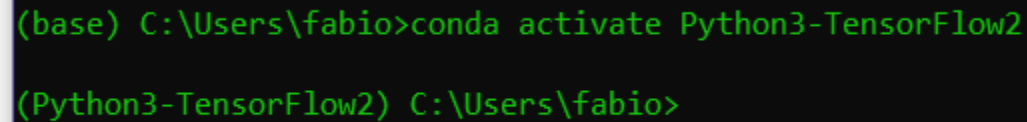
ca-certificates	pkgs/main/win-64::ca-certificates-2019.11.27-0
certifi	pkgs/main/win-64::certifi-2019.11.28-py38_0
openssl	pkgs/main/win-64::openssl-1.1.1d-he774522_3
pip	pkgs/main/win-64::pip-19.3.1-py38_0
python	pkgs/main/win-64::python-3.8.1-h5fd99cc_1
setuptools	pkgs/main/win-64::setuptools-44.0.0-py38_0
sqlite	pkgs/main/win-64::sqlite-3.30.1-he774522_0
vc	pkgs/main/win-64::vc-14.1-h0510ff6_4
vs2015_runtime	pkgs/main/win-64::vs2015_runtime-14.16.27012-hf0eaf9b_1
wheel	pkgs/main/win-64::wheel-0.33.6-py38_0
wincertstore	pkgs/main/win-64::wincertstore-0.2-py38_0

Proceed ([y]/n)? y

Installing TensorFlow 2.1.0

- Commands:

```
conda activate env_name
```

A terminal window with a black background and green text. The first line shows the command 'conda activate Python3-TensorFlow2' being executed from the base environment. The second line shows the prompt changing to '(Python3-TensorFlow2) C:\Users\fabio>' after successful activation.

```
(base) C:\Users\fabio>conda activate Python3-TensorFlow2  
(Python3-TensorFlow2) C:\Users\fabio>
```

```
conda install tensorflow
```

```
pip install --upgrade tensorflow
```

```
pip install ipykernel
```

ipykernel installation ensures that Jupyter has the right kernel

Installing TensorFlow 2.1.0

Finally:

```
python -m ipykernel install --name env_name
```

```
(Python3-TensorFlow2) C:\Users\fabio>python -m ipykernel install --name Python3-TensorFlow2  
Installed kernelspec Python3-TensorFlow2 in C:\ProgramData\jupyter\kernels\python3-tensorflow2
```

```
conda info --envs
```

```
(Python3-TensorFlow2) C:\Users\fabio>conda info --envs  
# conda environments:  
#  
base                C:\Users\fabio\Anaconda3  
Python3-TensorFlow2 * C:\Users\fabio\Anaconda3\envs\Python3-TensorFlow2
```

Installing TensorFlow 2.1.0

- Now, close and open again Jupyter to test if everything went right
- Be sure to select the right Kernel

```
In [2]: import tensorflow as tf  
        print(tf.__version__)
```

```
2.1.0
```

```
In [ ]:
```

Installing Packages

- Let's install scikit-learn and the TensorFlow datasets

```
pip install scikit-learn
```

```
pip install tensorflow-datasets
```


Let's “machine learning” with
Numpy

Elements of the model in supervised learning

- Inputs
- Weights
- Biases
- Targets
- Outputs

Import the relevant libraries

First: `pip install matplotlib`

Then:

```
In [2]: import numpy as np  
import matplotlib.pyplot as plt  
from mpl_toolkits.mplot3d import Axes3D
```

```
In [ ]:
```

Matplotlib and Axes3D are not required, but they make cool graphs

Generating random input data for training

We will start with a textbook example

- We generate random data with a linear relationship

```
In [3]: observations_nr = 1000
x_values = np.random.uniform(low=-10,high=10,size=(observations_nr, 1))
z_values = np.random.uniform(low=-10,high=10,size=(observations_nr, 1))

inputs = np.column_stack((x_values, z_values))
```

- We select how many samples/observations we need (`observations_nr`)
- We use `np.random.uniform` to generate these values
- We use `np.column_stack` to literally stack two vectors in a matrix
- Note that `inputs` is (1000 x 2)

Create Targets

We now define the linear function that will act as target for our model

For example: $f(x, z) = 4x - 3z + 2 + \text{noise}$

$\begin{matrix} \uparrow & \uparrow & \uparrow \\ w_1 & w_2 & b \end{matrix}$

- The **noise** ensures that the data looks more random
In fact, real data always contains noise

```
In [4]: noise = np.random.uniform(-1,1,(observations_nr,1))
        targets = 4*x_values - 3*z_values + 2 + noise
```


Let's set up the initial variables

We saw previously that for gradient descent we could select a random starting point

- In this case, though, it is better to force the hand a little bit

➤ We select **random small initial weights and biases**

We keep them within a small range



```
In [5]: boundary_range = 0.1
weights = np.random.uniform(-boundary_range, boundary_range, (2,1))
biases = np.random.uniform(-boundary_range, boundary_range, 1)
```

Note that \mathbf{W} is (2×1) , while \mathbf{b} is (1×1)

Now, let's set up the learning rate

- You can try different values and see what happens.
Here an example:

```
learning_rate = 0.02
```

Okay, we are all set.

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
from mpl_toolkits.mplot3d import Axes3D
```

```
In [2]: observations_nr = 1000
x_values = np.random.uniform(low=-10,high=10,size=(observations_nr, 1))
z_values = np.random.uniform(low=-10,high=10,size=(observations_nr, 1))

inputs = np.column_stack((x_values, z_values))
```

```
In [3]: noise = np.random.uniform(-1,1,(observations_nr,1))
targets = 4*x_values - 3*z_values + 2 + noise
```


Elements of the model in supervised learning

- ✓ Inputs
- ✓ Weights
- ✓ Biases
- ✓ Targets
- Outputs

We can now train the model

```
In [10]: for i in range (100):  
          outputs = np.dot(inputs, weights) + biases  
          deltas = outputs - targets  
  
          loss = np.sum(deltas**2) / 2 / observations_nr  
          print(loss)  
  
          deltas_scaled = deltas / observations_nr  
  
          weights = weights - learning_rate * np.dot(inputs.T, deltas_scaled)  
          biases = biases - learning_rate * np.sum(deltas_scaled)
```

- We run the algorithm over 100 iterations
This is an arbitrary number
- For this problem, it is more than enough


We can now train the model

```
In [10]: for i in range(100):  
         → outputs = np.dot(inputs, weights) + biases  
           deltas = outputs - targets  
  
           loss = np.sum(deltas**2) / 2 / observations_nr  
           print(loss)  
  
           deltas_scaled = deltas / observations_nr  
  
           weights = weights - learning_rate * np.dot(inputs.T, deltas_scaled)  
           biases = biases - learning_rate * np.sum(deltas_scaled)
```

- We calculate the outputs for the given weights and biases
They were random, so likely far from the targets

We can now train the model

```
In [10]: for i in range (100):  
         → outputs = np.dot(inputs, weights) + biases  
           deltas = outputs - targets  
  
           loss = np.sum(deltas**2) / 2 / observations_nr  
           print(loss)  
  
           deltas_scaled = deltas / observations_nr  
  
           weights = weights - learning_rate * np.dot(inputs.T, deltas_scaled)  
           biases = biases - learning_rate * np.sum(deltas_scaled)
```

$$f(x) = x_1 * w_1 + x_2 * w_2 + b$$

$$f(x) = \begin{matrix} & x \\ \begin{bmatrix} x_1 & x_2 \end{bmatrix} & \cdot \end{matrix} \begin{matrix} w \\ \begin{bmatrix} w_1 \\ w_2 \end{bmatrix} \end{matrix} + \begin{bmatrix} b \end{bmatrix}$$

We can now train the model

```
In [10]: for i in range (100):  
          outputs = np.dot(inputs, weights) + biases  
          → deltas = outputs - targets  
  
          loss = np.sum(deltas**2) / 2 / observations_nr  
          print(loss)  
  
          deltas_scaled = deltas / observations_nr  
  
          weights = weights - learning_rate * np.dot(inputs.T, deltas_scaled)  
          biases = biases - learning_rate * np.sum(deltas_scaled)
```

- We compute the deltas, that is, the error between outputs and targets

We can now train the model

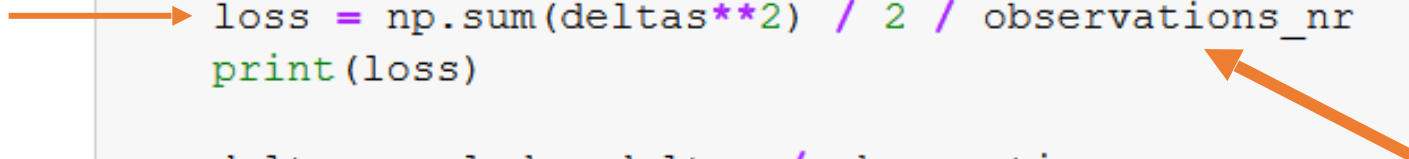
```
In [10]: for i in range (100):  
          outputs = np.dot(inputs, weights) + biases  
          deltas = outputs - targets  
  
          → loss = np.sum(deltas**2) / 2 / observations_nr  
             print(loss)  
  
          deltas_scaled = deltas / observations_nr  
  
          weights = weights - learning_rate * np.dot(inputs.T, deltas_scaled)  
          biases = biases - learning_rate * np.sum(deltas_scaled)
```

- We compute the Loss function that compares the outputs with the targets
- We used:

$$L(y, t) = \frac{L2 - norm}{2} = \frac{\sum_i (y_i - t_i)^2}{2}$$

We can now train the model

```
In [10]: for i in range(100):  
          outputs = np.dot(inputs, weights) + biases  
          deltas = outputs - targets  
  
          → loss = np.sum(deltas**2) / 2 / observations_nr  
             print(loss)  
  
          deltas_scaled = deltas / observations_nr  
  
          weights = weights - learning_rate * np.dot(inputs.T, deltas_scaled)  
          biases = biases - learning_rate * np.sum(deltas_scaled)
```



- We also **divided by the number of observations**
 - We do this to make the **learning independent of the number of observations**
Also, it does not affect the logic of the Loss (it's just a division by a constant)

We can now train the model

```
In [10]: for i in range (100):  
          outputs = np.dot(inputs, weights) + biases  
          deltas = outputs - targets  
  
          loss = np.sum(deltas**2) / 2 / observations_nr  
          → print(loss)  
  
          deltas_scaled = deltas / observations_nr  
  
          weights = weights - learning_rate * np.dot(inputs.T, deltas_scaled)  
          biases = biases - learning_rate * np.sum(deltas_scaled)
```

- We **print** the **loss** because we want to see if it is **decreasing**
- Otherwise, we need to change the **learning rate**

We can now train the model

```
In [10]: for i in range (100):  
          outputs = np.dot(inputs, weights) + biases  
          deltas = outputs - targets  
  
          loss = np.sum(deltas**2) / 2 / observations_nr  
          print(loss)  
  
          → deltas_scaled = deltas / observations_nr  
  
          weights = weights - learning_rate * np.dot(inputs.T, deltas_scaled)  
          biases = biases - learning_rate * np.sum(deltas_scaled)
```

- Even this step is done to make the algorithm independent of the number of observations


We can now train the model

```
In [10]: for i in range(100):
          outputs = np.dot(inputs, weights) + biases
          deltas = outputs - targets

          loss = np.sum(deltas**2) / 2 / observations_nr
          print(loss)

          deltas_scaled = deltas / observations_nr

          weights = weights - learning_rate * np.dot(inputs.T, deltas_scaled)
          biases = biases - learning_rate * np.sum(deltas_scaled)
```



- We update the weights and the biases following the gradient descent methodology

$$w_{i+1} = w_i - \eta \nabla_w L(y, t)$$

- Note the `inputs.T`

We are computing the transpose

$$b_{i+1} = b_i - \eta \nabla_b L(y, t)$$


We can now train the model

```
In [10]: for i in range(100):
          outputs = np.dot(inputs, weights) + biases
          deltas = outputs - targets

          loss = np.sum(deltas**2) / 2 / observations_nr
          print(loss)

          deltas_scaled = deltas / observations_nr

          weights = weights - learning_rate * np.dot(inputs.T, deltas_scaled)
          biases = biases - learning_rate * np.sum(deltas_scaled)
```



- We update the weights and the biases following the gradient descent methodology

$$w_{i+1} = w_i - \eta \sum_i x_i \delta_i$$

- Note the `inputs.T`

We are computing the transpose

$$b_{i+1} = b_i - \eta \sum_i \delta_i$$

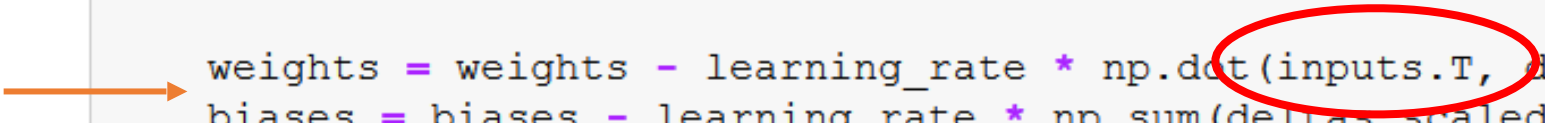
We can now train the model

```
In [10]: for i in range(100):
          outputs = np.dot(inputs, weights) + biases
          deltas = outputs - targets

          loss = np.sum(deltas**2) / 2 / observations_nr
          print(loss)

          deltas_scaled = deltas / observations_nr

          weights = weights - learning_rate * np.dot(inputs.T, deltas_scaled)
          biases = biases - learning_rate * np.sum(deltas_scaled)
```



- The transpose is required, because sometime `inputs` and `deltas_scaled` matrices cannot be multiplied (dot) together
In fact, it would be a $(1000 \times 2) \cdot (1000 \times 1)$
- However, `inputs.T` is (2×1000) and the dot product would now work

230.4481472811235
39.374264845768636
14.84691947374207
11.363488990284926
10.553401120641977
10.095731041651907
9.696665553008133
9.318518951803059
8.95599328186409
8.607905032668269
8.273611086149172
7.952556174035672
7.6442149018383745
7.348083677036639
7.063678974917241
6.790536433824226
6.528210082403527
6.276271609604928
6.034309665128121
5.80192918776475
5.578750760362001
5.364409990332156
5.158556914693805
4.960855428672885
4.770982736930607

...

0.7525360163979564
0.7293088446293438
0.7070014471786221
0.6855774018414768
0.665001728698878
0.6452408330039354
0.6262624503303736
0.6080355938930729
0.590530503954679
0.5737185992356539
0.5575724302484455
0.5420656344795837
0.5271728933465191
0.5128698908589369
0.49913327391703893
0.4859406141819728
0.47327037145617024
0.46110185851377616
0.44941520732376616
0.43819133661059567
0.42741192069942135
0.4170593595950113
0.40711675024551575
0.3975678589441579
0.3883970948237942

It's fast

It's minimizing the error (Loss function)

We trained the model correctly!

Bonus: let's check

```
print(weights, biases)
```

```
[[ 3.99920143]  
 [-2.99934675]] [1.75006155]
```

$$f(x, z) = 4x - 3z + 2 + \text{noise}$$

The weights seem right.

The bias... almost. Why is that?

- Probably, not enough iterations (100) or not perfect learning rate

Bonus: let's check

```
print(weights, biases)
```

```
[[ 3.99920143]  
 [-2.99934675]] [1.75006155]
```

$$f(x, z) = 4x - 3z + 2 + \text{noise}$$

- Let's run again the block in Jupyter that contains the loop
- This will run just that part of the code, so obtaining other 100 iterations

Bonus: let's check

```
print(weights, biases)
```

```
[[ 3.99645043  
  -2.99907267]] [1.98452474]
```

$$f(x, z) = 4x - 3z + 2 + \text{noise}$$

If we print again the values we get so much closer!

Recap of the “tuning” parameters

- Number of observations
- Learning rate
- Number of iterations (or, better, **epochs**)
- Initial range for initializing the weights and biases